

ECE 406 COURSE NOTES
ALGORITHM DESIGN AND ANALYSIS

Paolo Torres

University of Waterloo
Winter 2021

Table of Contents

1	<i>Introduction and Basic Arithmetic</i>	I
1.1	Algorithms, Correctness, Termination, Efficiency	1
1.1.1	Algorithms.....	1
1.1.2	Correctness	1
1.1.3	Termination.....	1
1.1.4	Time Efficiency	2
1.1.5	Claim 1.....	2
1.1.6	More Efficient Algorithm	3
1.1.7	Note on Measuring Time Efficiency	3
1.2	Big-O Notation	3
1.2.1	Definition 1 (O).....	3
1.2.2	Definition 2 (Ω)	3
1.2.3	Definition 3 (Θ)	3
1.2.4	Example.....	4
1.2.5	Big-O Explanation.....	4
1.2.6	Big-O Simplifications	5
1.3	Arithmetic	5
1.3.1	Addition.....	5
1.3.2	Multiplication.....	5
1.3.3	Division.....	6
2	<i>Algorithms with Numbers</i>	8
2.1	Modular Arithmetic	8
2.1.1	Example Application: Two's Complement Arithmetic	8
2.1.2	Modular Addition, Subtraction	8
2.1.3	Modular Multiplication	9
2.1.4	Modular Exponentiation	9

2.1.5	Towards Modular Division: GCD Using Euclid	9
2.1.6	Towards Modular Division: Extended Euclid	10
2.1.7	Modular Division.....	11
2.2	Primality Testing.....	11
2.2.1	Fermat's Little Theorem.....	12
2.3	Generating an n-Bit Prime.....	13
2.4	Cryptography.....	14
2.4.1	Claim 1.....	14
2.4.2	Example.....	14
2.4.3	Attacks.....	15
2.4.4	Proof for Claim 1	15
3	<i>Universal Hashing, Divide and Conquer.....</i>	17
3.1	Universal Hashing	17
3.1.1	Example.....	17
3.1.2	Two Possible Candidates for \mathcal{D}	17
3.1.3	Hash Table Objectives.....	18
3.1.4	Problem Setup	18
3.1.5	How a Hash Table Works	18
3.1.6	Universal Hashing	19
3.1.7	How to Construct \mathcal{H} Given U, m	20
3.2	Divide and Conquer	22
3.2.1	Definition	22
3.2.2	Solving Recurrences	24
3.2.3	Sorting Problem	25
3.3	Divide and Conquer Applications.....	26
3.3.1	Selection Problem	26
3.3.2	Matrix Multiplication	28

4	<i>Decompositions of Graphs, Paths</i>	30
4.1	Graphs.....	30
4.1.1	Representing a Graph	30
4.1.2	Graph Exploration	31
4.2	DFS, Continued	33
4.2.1	Directed Acyclic Graph (DAG)	34
4.2.2	Strongly Connected Components	34
4.3	Shortest Paths	37
4.3.1	Breadth First Search (BFS).....	37
5	<i>Paths in Graphs, Continued</i>	40
5.1	Shortest Paths in Weighted Graphs	40
5.1.1	Dijkstra's Algorithm.....	40
5.1.2	Correctness	43
5.1.3	Time Efficiency of Dijkstra.....	45
5.2	Shortest Paths and Negative Edge Weights.....	47
5.2.1	Bellman-Ford Algorithm.....	48
5.2.2	Directed Acyclic Graphs (DAGs)	49

1 INTRODUCTION AND BASIC ARITHMETIC

1.1 Algorithms, Correctness, Termination, Efficiency

1.1.1 Algorithms

Given the specification for a function, an algorithm is the procedure to compute it.

Example:

$$F: \mathbb{Z}_o^+ \rightarrow \mathbb{Z}_o^+, \text{ where } F(n) = \begin{cases} 0, & \text{if } n = 0 \\ 1, & \text{if } n = 1 \\ F(n - 1) + F(n - 2), & \text{otherwise} \end{cases}$$

Commonly used sets: N, Z (all ints), Z^+ (positive ints), \mathbb{Z}_o^+ (non-negative ints), R, \dots

Fibonacci Sequence:

```
FIB1(n)
1 if n = 0 then return 0
2 if n = 1 then return 1
3 return FIB1(n - 1) + FIB1(n - 2)
```

Important aspects:

- Function has been specified as a recurrence, so a recursive algorithm seems natural
- Imperative (procedural) specification of an algorithm has consequences:
 - Intuiting correctness can be a challenge
 - Intuiting time and space efficiency may be easier
- No mundane error checking, can focus on core logic
- Input value n is unbounded but finite

1.1.2 Correctness

Correctness refers to an algorithm's ability to guarantee expected termination. In the case of FIB_1 , it is a direct encoding of the recurrence.

1.1.3 Termination

The end of an algorithm. It can be proven that FIB_1 terminates on every input $n \in \mathbb{Z}_o^+$ by induction on n .

1.1.4 Time Efficiency

Can be calculated by counting the number of: (i) comparisons – these happen on Lines (1) and (2), and (ii) number of additions – this happens on Line (3).

Suppose $T(n)$ represents the time efficiency of FIB_1 :

$$T(n) = \begin{cases} 1, & \text{if } n = 0 \\ 2, & \text{if } n = 1 \\ 3 + T(n - 1) + T(n - 2), & \text{otherwise} \end{cases}$$

How bad is $T(n)$? Is it exponential in n ?

For all n , $T(n) \geq F(n)$.

1.1.5 Claim 1

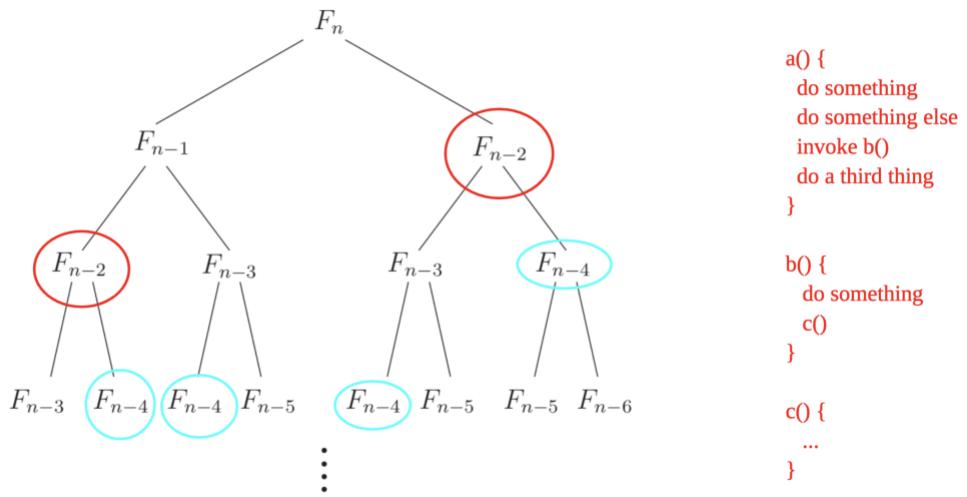
For all $n \in \mathbb{Z}_o^+$, $F(n) \geq (\sqrt{2})^n$.

If this claim is true, then $T(n) \geq (\sqrt{2})^n$, and because $\sqrt{2} > 1$, $T(n)$ is exponential in n .

Proof for the claim: by induction on n .

Does a better algorithm exist from the standpoint of time efficiency?

Figure 0.1 The proliferation of recursive calls in `fib1`.



Recall how subroutine (recursive, in this case) invocation works:

- Every node in the tree corresponds to an invocation of the algorithm
- Sequence of invocations corresponds to a pre-order traversal
- Maximum depth of the call stack at any moment: n

Main point in this case: Redundancy, F_i , appears more than once.

1.1.6 More Efficient Algorithm

```
FIB2(n)
1 if n = 0 then return 0
2 create an array f[0, ..., n]
3 f[0] ← 0, f[1] ← 1
4 foreach i from 2 to n do
5   f[i] ← f[i - 1] + f[i - 2]
6 return f[n]
```

Let $U(n)$ be the # of comparisons plus additions on input n :

$$U(n) = \begin{cases} 1, & \text{if } n = 0 \\ n, & \text{otherwise} \end{cases}$$

Linear in n for $n \geq 1$, more efficient than FIB_1 .

1.1.7 Note on Measuring Time Efficiency

Need to pick the right level of abstraction, meaning picking some kind of “hot spot” or “hot operation,” then count. For example, number of additions, comparisons, recursive calls, etc.

1.2 Big-O Notation

1.2.1 Definition 1 (O)

Let $f: N \rightarrow R^+$, and $g: N \rightarrow R^+$ be functions. Define $f = O(g)$ if there exists a constant $c \in R^+$ such that $f(n) \leq c \cdot g(n)$.

- $N = Z^+: \{1, 2, 3, \dots\}, R^+: \text{set of positive real numbers}$
- Typically consider non-decreasing functions only

1.2.2 Definition 2 (Ω)

Define $f = \Omega(g)$ if $g = O(f)$

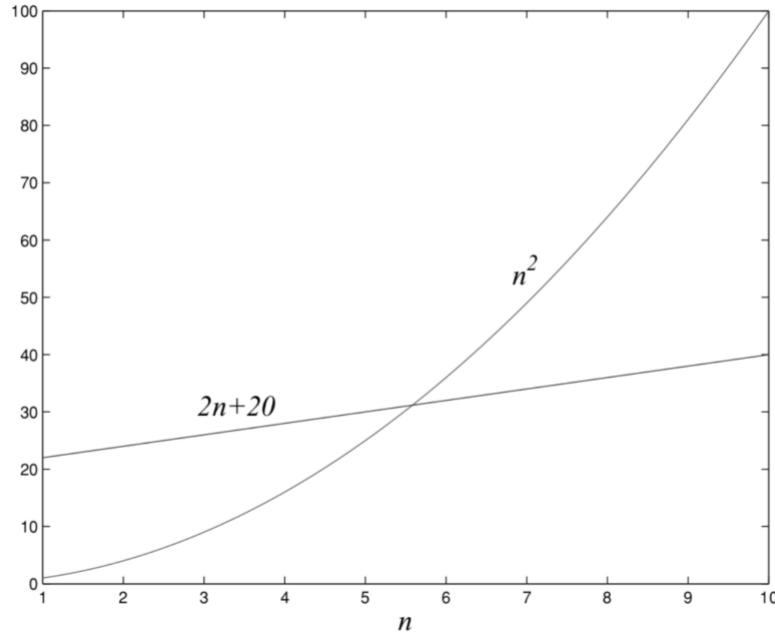
1.2.3 Definition 3 (Θ)

Define $f = \Theta(g)$ if $f = O(g)$ and $f = \Omega(g)$

- $f = O(g)$ analogous to $f \leq g$
- $f = \Omega(g)$ analogous to $f \geq g$
- $f = \Theta(g)$ analogous to $f = g$

1.2.4 Example

Figure 0.2 Which running time is better?



Precise answer to this question: depends on n .

But in big-O notation:

- $2n + 20 = O(n^2)$. Proof: Adopt as the constant $c \in R$ for any $c > 22$
- $2n + 2 \neq \Omega(n^2) \therefore 2n + 2 \neq \Theta(n^2)$

1.2.5 Big-O Explanation

Suppose algorithm A runs in $2n + 20$ time, B in n^2 , and C in 2^n . Now suppose the speed of the computer doubles, which algorithm gives the best payoff?

For a given time period t , what is the largest input n each algorithm can handle? Set $runtime = t$ and $runtime = 2t$, solve for n :

Algorithm	Old Computer	New Computer
A	$t/2 - 10$	$t - 10$
B	\sqrt{t}	$\sqrt{2} \cdot \sqrt{t}$
C	$\log_2 t$	$1 + \log_2 t$

So, payoff with algorithm A is approximately $2x$, B is $1.4x$, and C is $1 +$.

1.2.6 Big-O Simplifications

- Multiplicative and additive constants can be omitted
- n^a dominates n^b for $a > b \geq 0$
- Any exponential dominates any polynomial, any polynomial dominates any logarithm
- Big-O simplifications should be used prudently, not applicable in all settings

1.3 Arithmetic

1.3.1 Addition

Hypothesize access to a function $T : \{0, 1\} \times \{0, \dots, 9\} \times \{0, \dots, 9\} \rightarrow \{0, 1\} \times \{0, \dots, 9\}$:

Carry	One Digit	Other Digit	Result Carry	Result Sum
0	0	0	0	0
1	0	0	0	1
0	0	1	0	1
...				
1	8	9	1	8
0	9	9	1	8
1	9	9	1	9

To add 7,814 and 93,404:

$$\begin{aligned} C &\rightarrow 111000 \\ 1 &\rightarrow 007814 \\ 2 &\rightarrow 093404 \\ T &\rightarrow 101218 \end{aligned}$$

digits needed to encode $x \in \mathbb{Z}^+ = \lceil \log_{10} x \rceil + 1$

bits needed to encode $x \in \mathbb{Z}^+ = \lceil \log_2 x \rceil + 1$

So, time efficiency of an algorithm to add $x, y \in \mathbb{Z}^+$ as measured by number of lookups to T:

- $1 + \lceil \log_{10}(\max\{x, y\}) \rceil$ in the best case
- $2 + \lceil \log_{10}(\max\{x, y\}) \rceil$ in the worst case
- So, either way, $\Theta(n)$, or linear time, where n is the size of the input

1.3.2 Multiplication

For $x, y \in \mathbb{Z}_o^+$, encoded in binary:

$$x \times y = \begin{cases} 0, & \text{if } y = 0 \\ 2(x \times \lfloor y/2 \rfloor), & \text{if } y \text{ even, } y > 0 \\ x + 2(x \times \lfloor y/2 \rfloor), & \text{otherwise} \end{cases}$$

Straightforward encoding as recursive algorithm $MULTIPLY(x, y)$.

Figure 1.1 Multiplication à la Français.

```
function multiply(x, y)
```

Input: Two n -bit integers x and y , where $y \geq 0$

Output: Their product

```
if  $y = 0$ : return 0
z = multiply( $x, \lfloor y/2 \rfloor$ )
if  $y$  is even:
    return  $2z$ 
else:
    return  $x + 2z$ 
```

Worst case running time:

- Let # bits to encode each of x and $y = n$
- # recursive calls = $\Theta(n)$
- In each call:
 - One comparison to 0, one division by 2 (right bit shift), one assignment to z , one check for evenness (check LSB), one multiplication by 2 (left bit shift), one addition of $O(n)$ -bit numbers
 - So, $O(n^2)$ in the worst case

1.3.3 Division

Definition 1: Given $x \in Z_o^+, y \in Z^+$, the pair $\langle q, r \rangle$ where $q \in Z_o^+, r \in \{0, 1, \dots, y - 1\}$ of x divided by y are those that satisfy:

$$x = q \cdot y + r$$

Claim 1: For every $x \in Z_o^+, y \in Z^+$, $\langle q, r \rangle$ as defined above (i) exists, and (ii) is unique.

To specify a recurrence for $\langle q, r \rangle$, denote as $\langle q', r' \rangle$, the result of $\lfloor x/2 \rfloor$ divided by y . Now:

$$\langle q, r \rangle = \begin{cases} \langle 0, 0 \rangle, & \text{if } x = 0 \\ \langle 2q', 2r' \rangle, & \text{if } x \text{ even and } 2r' < y \\ \langle 2q', 2r' + 1 \rangle, & \text{if } x \text{ odd and } 2r' + 1 < y \\ \langle 2q' + 1, 2r' - y \rangle, & \text{if } x \text{ even and } 2r' \geq y \\ \langle 2q' + 1, 2r' + 1 - y \rangle, & \text{otherwise} \end{cases}$$

Claim 2: The above recurrence is correct.

Proof: Cases are exhaustive. Proof by case-analysis and induction on # bits to encode x .

By induction assumption: $0 \leq r' \leq y - 1 \therefore 0 \leq 2r' \leq 2y - 2$.

Figure 1.2 Division.

function divide(x, y)

Input: Two n -bit integers x and y , where $y \geq 1$

Output: The quotient and remainder of x divided by y

```
if  $x = 0$ : return  $(q, r) = (0, 0)$ 
 $(q, r) = \text{divide}(\lfloor x/2 \rfloor, y)$ 
 $q = 2 \cdot q$ ,  $r = 2 \cdot r$ 
if  $x$  is odd:  $r = r + 1$ 
if  $r \geq y$ :  $r = r - y$ ,  $q = q + 1$ 
return  $(q, r)$ 
```

Running time: $O(n^2)$.

2 ALGORITHMS WITH NUMBERS

2.1 Modular Arithmetic

- **Definition 1:** For $x \in Z, N \in Z^+, x \text{ modulo } N$ is the remainder of x divided by N
- **Definition 2:** $x \equiv y \pmod{N}$ if N divides $x - y$. " \equiv " read as "congruent to"
 - Example: $373 \equiv 13 \pmod{60}, 59 \equiv -1 \pmod{60}$

2.1.1 Example Application: Two's Complement Arithmetic

Suppose we want to represent, using n bits, positive and negative integers, and 0. Could reserve 1 bit for sign. This would allow us to represent integers in the interval $[-(2^{n-1} - 1), 2^{n-1} - 1]$, with a "positive zero" and a "negative zero."

In two's complement arithmetic, we have exactly one bit-string for 0, and represent integers in the interval $[-2^{n-1}, 2^{n-1} - 1]$. How? Represent any $x \in [-2^{n-1}, 2^{n-1} - 1] \cap Z$ as the non-negative integer modulo 2^n . So:

$$0 \leq x \leq 2^{n-1} - 1 \rightarrow x \text{ is represented as } x$$

- Ex: For $n = 5, (9)_{10}$ written as $(01001)_2$
 $-2^{n-1} \leq x < 0 \rightarrow x \text{ is represented as } 2^n + x$
- Ex: For $n = 5, (-9)_{10} \equiv 32 + (-9) \equiv 23 \pmod{32}$, written as $(10111)_2$

All arithmetic performed modulo 2^n :

- Ex: For $n = 5, 13 + (-7) \equiv 01101 + 11001 \equiv 100110 \equiv 00110 \equiv 6 \pmod{32}$

Claim 1: $x \equiv x', y \equiv y' \pmod{N}$ implies: $x + y \equiv x' + y', xy \equiv x'y' \pmod{N}$

Claim 2:

$$\begin{aligned} x + (y + z) &\equiv (x + y) + z \pmod{N} \\ xy &\equiv yx \pmod{N} \\ x(y + z) &\equiv xy + xz \pmod{N} \end{aligned}$$

Example: $2^{3045} \equiv (2^5)^{609} \equiv (1)^{609} \equiv 1 \pmod{31}$

2.1.2 Modular Addition, Subtraction

$$x + y \equiv \begin{cases} x + y \pmod{N}, & \text{if } 0 \leq x + y < N \\ x + y - N \pmod{N}, & \text{otherwise} \end{cases}$$

$$x - y \equiv \begin{cases} x - y \pmod{N}, & \text{if } 0 \leq x - y < N \\ x - y + N \pmod{N}, & \text{otherwise} \end{cases}$$

- Any intermediate result is between $-(N - 1)$ and $2(N - 1)$
- So, time efficiency is $O(n)$, where $n = \lceil \log N \rceil$

2.1.3 Modular Multiplication

Let $MULT$ and DIV be our algorithms for non-modular multiplication and division:

$$x \times y \equiv r \pmod{N}, \text{ where } \langle q, r \rangle = DIV(MULT(x, y), N)$$

- Any intermediate result (specifically, result of $MULT$) is between 0 and $(N - 1)^2$
- So, time efficiency is $O(n^2)$, where $n = \lceil \log N \rceil$

2.1.4 Modular Exponentiation

- Recall: We used “repeated doubling” for non-modular multiplication and “repeated halving” for non-modular division
- Similarly, here, use “repeated squaring”:

$$x^y = \begin{cases} 1, & \text{if } y = 0 \\ (x^2)^{\lfloor y/2 \rfloor}, & \text{if } y \text{ is even} \\ x \cdot (x^2)^{\lfloor y/2 \rfloor}, & \text{otherwise} \end{cases}$$

Figure 1.4 Modular exponentiation.

```
function modexp(x, y, N)
Input: Two n-bit integers x and N, an integer exponent y
Output:  $x^y \pmod{N}$ 

if  $y = 0$ : return 1
z = modexp(x,  $\lfloor y/2 \rfloor$ , N)
if  $y$  is even:
    return  $z^2 \pmod{N}$ 
else:
    return  $x \cdot z^2 \pmod{N}$ 
```

Time Efficiency: $O(n^3)$

2.1.5 Towards Modular Division: GCD Using Euclid

- In a non-modular world, $a/b = a \times b^{-1}$. Only case for which b^{-1} doesn’t exist: $b = 0$
- In a modular world, $b^{-1} \pmod{N}$ may not exist even if $b \not\equiv 0 \pmod{N}$

Crucial building block: GCD. Example: What is $\gcd(1035, 759)$?

$$1035 = 3^2 \cdot 5 \cdot 23 \rightarrow 759 = 3 \cdot 11 \cdot 23 \rightarrow \gcd(1035, 759) = 3 \cdot 23$$

But factoring into prime factors conjectured to be computationally hard in the worse case

Claim 3: $x, y \in \mathbb{Z}^+, x \geq y \rightarrow \gcd(x, y) = \gcd(x \bmod y, y)$

Proof: Suffices to prove: $\gcd(x, y) = \gcd(x - y, y)$. Now prove \leq and \geq .

Figure 1.5 Euclid's algorithm for finding the greatest common divisor of two numbers.

```
function Euclid(a, b)
```

Input: Two integers a and b with $a \geq b \geq 0$

Output: $\gcd(a, b)$

```
if  $b = 0$ : return  $a$ 
```

```
return Euclid( $b, a \bmod b$ )
```

How fast does *Euclid* converge?

Claim 4: $a \geq b \geq 0 \rightarrow a \bmod b < a/2$

So: Guaranteed to lose at least 1 bit for every recursive call \rightarrow time efficiency is $O(n^3)$

2.1.6 Towards Modular Division: Extended Euclid

Claim 5: d divides a and b , and $d = ax + by$ for some $x, y \in \mathbb{Z} \rightarrow d = \gcd(a, b)$

Claim 6: Let $d = \gcd(a, b)$, $d = ax + by$ and $d = bx' +$

$(a \bmod b)y'$ for some $x, y, x', y' \in \mathbb{Z}$. Then:

$$\langle x, y \rangle = \begin{cases} \langle 1, 0 \rangle, & \text{if } b = 0 \\ \langle y', x' - \lfloor a/b \rfloor y' \rangle, & \text{otherwise} \end{cases}$$

Figure 1.6 A simple extension of Euclid's algorithm.

```
function extended-Euclid(a, b)
```

Input: Two positive integers a and b with $a \geq b \geq 0$

Output: Integers x, y, d such that $d = \gcd(a, b)$ and $ax + by = d$

```
if  $b = 0$ : return  $(1, 0, a)$ 
```

```
 $(x', y', d) = \text{extended-Euclid}(b, a \bmod b)$ 
```

```
return  $(y', x' - \lfloor a/b \rfloor y', d)$ 
```

Example run on input $\langle 359, 82 \rangle$:

Arguments	Return Value
$\langle 359, 82 \rangle$	$\langle -37, 162, 1 \rangle$
$\langle 82, 31 \rangle$	$\langle 14, -37, 1 \rangle$
$\langle 31, 20 \rangle$	$\langle -9, 14, 1 \rangle$
$\langle 20, 11 \rangle$	$\langle 5, -9, 1 \rangle$
$\langle 11, 9 \rangle$	$\langle -4, 5, 1 \rangle$
$\langle 9, 2 \rangle$	$\langle 1, -4, 1 \rangle$
$\langle 2, 1 \rangle$	$\langle 0, 1, 1 \rangle$
$\langle 1, 0 \rangle$	$\langle 1, 0, 1 \rangle$

A good example to segue to our final step in modular division.

We figured out: $\gcd(359, 82) = 1$, which implies:

- $82 \times 162 \equiv 1 \pmod{359}$. So: 162 is multiplicative inverse of 82 modulo 359
- So, for example: $116 \text{ divided by } 82 \text{ modulo } 359 \equiv 116 \times 162 = 124 \pmod{359}$

2.1.7 Modular Division

Definition 3: x is the multiplicative inverse of a modulo N if $ax \equiv 1 \pmod{N}$

Claim 7: For every $a \in \{0, \dots, N-1\}$, there exists at most $a^{-1} \pmod{N}$

Claim 8: Given $\langle a, N \rangle$, where $a \in \{0, \dots, N-1\}$, $a^{-1} \pmod{N}$ may not exist

Definition 4: If $\gcd(x, y) = 1$, then we say that x is relatively prime to y

Claim 9: $a^{-1} \pmod{N}$ exists if and only if a and N are relatively prime

So, to compute $a/b \pmod{N}$:

- Determine whether $\gcd(b, N) = 1$
- If yes to (i) determine $b^{-1} \pmod{N}$, and
- If yes to (i) compute $a \times b^{-1} \pmod{N}$

(i) and (ii) are done simultaneously by *extended – Euclid*

Running Time: (i) + (ii) = $O(n^3)$, (iii) = $O(n^2)$. So (i) + (ii) + (iii) = $O(n^3)$

2.2 Primality Testing

Given $n \in \mathbb{Z}^+$, is n prime?

For a decision problem, i.e., co-domain of function to be computed as $\{\text{true}, \text{false}\}$, a randomized algorithm:

- Has access to an unbiased coin
- Is deemed to be correct if:
 - $\Pr\{\text{Algorithm outputs false} \mid \text{input instance is false}\} = 1$
 - $\Pr\{\text{Algorithm outputs true} \mid \text{input instance is true}\} \geq 1/2$

Suppose:

- We run such an algorithm k times, pairwise independently
- We return *true* if and only if every run returns *true*
- Then, $\Pr\{\text{we return true incorrectly}\} \leq 2^{-k}$

2.2.1 Fermat's Little Theorem

Claim 1: p prime \rightarrow for all $a \in [1, p) \cap \mathbb{Z}$, $a^{p-1} \equiv 1 \pmod{p}$

To prove Fermat's little theorem, leverage the following:

Claim 2: p prime, $a, i, j \in \{1, 2, \dots, p-1\}$ and $i \neq j \rightarrow a \cdot i \not\equiv a \cdot j \pmod{p}$

Proof for Claim 2: We know that a, p are relatively prime, so $a^{-1} \pmod{p}$ exists.

$$a \cdot i \equiv a \cdot j \pmod{p} \rightarrow a \cdot i \cdot a^{-1} \equiv a \cdot j \cdot a^{-1} \pmod{p} \rightarrow i \equiv j \pmod{p} \rightarrow i = j$$

Proof for Claim 1: From Claim 2:

$$\{1, 2, \dots, p-1\} = \{a \cdot 1 \pmod{p}, a \cdot 2 \pmod{p}, \dots, a \cdot (p-1) \pmod{p}\}$$

Also, $(p-1)!^{-1} \pmod{p}$ exists.

So, $(p-1)! \equiv a^{p-1} \cdot (p-1)! \pmod{p} \rightarrow a^{p-1} \equiv 1 \pmod{p}$

Figure 1.7 An algorithm for testing primality.

```

function primality(N)
Input: Positive integer N
Output: yes/no

Pick a positive integer a < N at random
if  $a^{N-1} \equiv 1 \pmod{N}$ :
    return yes
else:
    return no

```

Issues with the algorithm:

1. Fermat's little theorem is not an “if and only if”: Carmichael numbers.

2. Suppose N is not prime/Carmichael. For a chosen, say, uniformly from $\{1, \dots, N - 1\}$, what is $\Pr\{a^{N-1} \not\equiv 1 \pmod{N}\}$?

- We know such an a exists, but how likely is it that we will pick it?

We do not deal with (1) – cop out: Carmichael numbers are rare.

Claim 3: If $a^{N-1} \not\equiv 1 \pmod{N}$ for a, N relatively prime, then it must hold for at least half the choices $a \in \{1, \dots, N - 1\}$.

Proof: If there exists no $b \in \{1, \dots, N - 1\}$ with $b^{N-1} \equiv 1 \pmod{N}$, then we are done.

If such a b exists, then $(b \cdot a)^{N-1} \not\equiv 1 \pmod{N}$.

Also, if b, c exist with $b \neq c$, $b^{N-1} \equiv c^{N-1} \equiv 1 \pmod{N}$, then:

$$b \cdot a \not\equiv c \cdot a \pmod{N}$$

So, at least as many $\not\equiv 1 \pmod{N}$ as there are $\equiv 1 \pmod{N}$.

So:

$$\Pr\{\text{Algorithm 1.7 returns yes when } N \text{ is prime/Carmichael}\} = 1$$

$$\Pr\{\text{Algorithm 1.7 returns yes when } N \text{ is not prime/Carmichael}\}, < 1/2$$

For k runs of Algorithm 1.7 on uniform, independent choices of a :

$$\Pr\{\text{Algorithm 1.7 returns yes on all } k \text{ runs when } N \text{ is not prime/Carmichael}\} \leq 2^{-k}$$

2.3 Generating an n-Bit Prime

Claim 4: $\Pr\{\text{uniformly chosen } n\text{-bit number is prime}\} \approx 1/n$.

So, algorithm for generating a prime:

1. Randomly generate n -bit number, r .
2. Check whether r is prime.
3. If not, go to Step (1).

Guaranteed return in Step (2) if r is indeed prime.

Each trial in the above algorithm is a Bernoulli trial:

- Only one of two outcomes: success or failure

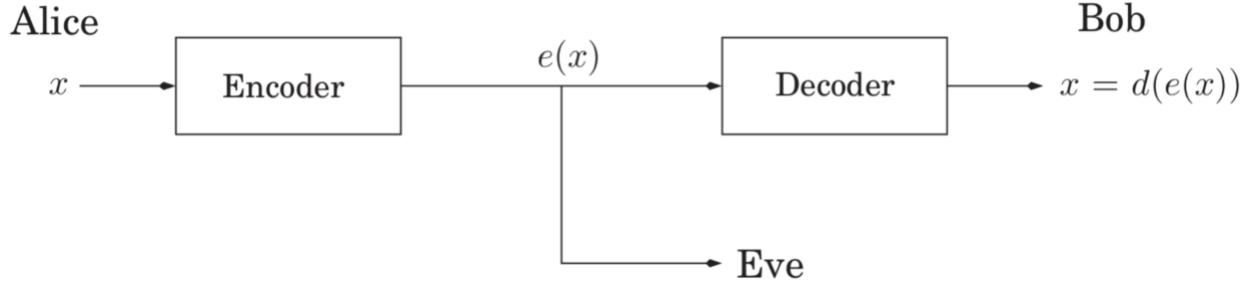
In a Bernoulli trial if $\Pr\{\text{success}\} = p$, then *expected # trials to see a success* = $1/p$

- Example: *Expected # tosses of fair coin to see a heads* = 2
- Example: *Expected # tosses of fair die to see, say, a 3* = 6

So, we expect the above algorithm to halt in n iterations.

2.4 Cryptography

RSA: Exploits presumed computational hardness of factoring vs. computational ease of GCD, primality testing and modular exponentiation.



Symmetric Key Cryptography:

- $e(\cdot) = d(\cdot) \rightarrow$ Alice and Bob both know $e(\cdot)$ and $d(\cdot)$
 - Example: $e(x) = x \oplus r, d(y) = y \oplus r$
- They keep these secrets from everyone else
- Bootstrapping Problem: How do Alice and Bob share $e(\cdot), d(\cdot)$?

Public Key Cryptography:

- Bob publishes $e(\cdot)$ to the whole world
- Bob keeps $d(\cdot)$ to himself
- RSA is an example of a public key cryptography scheme

2.4.1 Claim 1

Let p, q be primes and $N = pq$. For any e relatively prime to $(p - 1)(q - 1)$:

1. The function $f: \{0, 1, \dots, N - 1\} \rightarrow \{0, 1, \dots, N - 1\}$ where $f(x) = x^e \text{ mod } N$ is a bijection
2. Let $d = e^{-1} \text{ (mod } (p - 1)(q - 1)\text{)}.$ Then, $(x^e)^d \equiv x \text{ (mod } N\text{)}$

Bob publishes the pair $\langle N, e \rangle$. Alice encodes message x as $f(x)$ from the claim.

Bob keeps the d in the claim secret to themselves.

2.4.2 Example

$p = 11, q = 17$. Then, $N = pq = 187$. The only messages that can be sent: $\{0, 1, \dots, 186\}$.

We could pick $e = 7$: $\gcd(10 \times 16, 7) = 1$

Then, $d = 23$: $7^{-1} \equiv 23 \text{ (mod } 160\text{)}$.

To send the message 98, Alice would send $98^7 \bmod 187 = 21$.

Bob would decode the message as $21^{23} \bmod 187 = 98$.

2.4.3 Attacks

Attacker knows: (i) $\langle N, e \rangle$, (ii) $x^e \bmod N$.

- Attack 1: Attacker determines x given (i) and (ii).
- Attack 2: Even more devastating – attacker factors $N = p \cdot q$.
 - They can then compute $d = e^{-1} \pmod{(p-1)(q-1)}$.

2.4.4 Proof for Claim 1

Property #2 implies Property #1.

To prove Property #2:

Because $ed \equiv 1 \pmod{(p-1)(q-1)}$, $ed = 1 + k(p-1)(q-1)$ for some $k \in \mathbb{Z}$.

We seek to show: $x^{ed} - x = x^{1+k(p-1)(q-1)} - x$ is divisible by N , and is therefore $\equiv 0 \pmod{N}$.

Now, by Fermat's little theorem: $x \cdot (x^{p-1})^{k(q-1)} - x \equiv 0 \pmod{p}$.

And again, by Fermat's little theorem: $x \cdot (x^{q-1})^{k(p-1)} - x \equiv 0 \pmod{q}$.

So: $x^{ed} - x$ is divisible by the product of the two primes $pq = N$.

Figure 1.9 RSA.

-
- Bob chooses his public and secret keys.
- He starts by picking two large (n -bit) random primes p and q .
 - His public key is (N, e) where $N = pq$ and e is a $2n$ -bit number relatively prime to $(p - 1)(q - 1)$. A common choice is $e = 3$ because it permits fast encoding.
 - His secret key is d , the inverse of e modulo $(p - 1)(q - 1)$, computed using the extended Euclid algorithm.
- Alice wishes to send message x to Bob.
- She looks up his public key (N, e) and sends him $y = (x^e \bmod N)$, computed using an efficient modular exponentiation algorithm.
 - He decodes the message by computing $y^d \bmod N$.
-

Example:

```
$ ssh-keygen -t rsa
...
$ cat id_rsa.pub
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQDJAs5HIayjHG
LdvEeiaRI2R3TG8+chfGYrjEWc82bV3ndC87+dYAFVXyVDdc2C
0vDHY6cNcN4vpjc0fZbieeJWC0wjFV8qt5VZDTdvtLJSBi1H1j
1JI6FoGBjwMjqoDsXR0n3e7rundqr0xLsk6RoIVunhluloj2Ss
L2fwU7/pbhrvWBBx1jP6aaCkW5sAEu143xM71C2bAMqzoS47WY
+xH91sgm8hwji/KUEoHeeVMrc54bTsozPQp4t+3QwjDfqEMeTy
BvJ93ZTZFHJiQVORTnw3x8HyzNgYTDZVnFJi6kWx8suggcokgz
ffAM+7xNJ1zDmJ1bY3N+WHuRDsvFpf me@localhost
```

3 UNIVERSAL HASHING, DIVIDE AND CONQUER

3.1 Universal Hashing

3.1.1 Example

University of Waterloo has data about students that looks something like:

Student ID	Last Name	First Name	Degree Objective	Date of Matriculation	Term 1 + Marks	...
------------	-----------	------------	------------------	-----------------------	----------------	-----

The Student ID is used to uniquely identify each student record. We call it a *key*. The remainder of the data for each student is called *satellite data*. The school seeks to store this data and retrieve it efficiently. Assume that \mathcal{D} is the data structure that stores all the data. For any given student record, r , there is an efficient function $key(r)$, which returns the key, i.e., Student ID, associated with r .

- $CreateNew(\dots)$: Creates a new, empty instance of \mathcal{D} and returns it. May take arguments, e.g., the space allocated to \mathcal{D} .
- $Insert(\mathcal{D}, r)$: Insert the record r into \mathcal{D} . If a record with $key(r)$ already exists in \mathcal{D} , adopt some convention, e.g., replace it with this new one.
- $Delete(\mathcal{D}, k)$: Remove the record associated with key k in \mathcal{D} . If such a record does not exist, adopt some convention, e.g., leave \mathcal{D} unchanged.
- $Search(\mathcal{D}, k)$: Return the record associated with key k from \mathcal{D} . If such a record does not exist, return a special string, e.g., ϵ .

What data structure, and associated algorithms, will you choose for \mathcal{D} ?

Answer: It depends on: (i) how frequently *Search* is invoked compared to *Insert* and/or *Delete*, (ii) how much space \mathcal{D} is allowed to take, (iii) how efficiently *Search* needs to run.

3.1.2 Two Possible Candidates for \mathcal{D}

Array, indexed by Student ID:

- Good: *Search* runs in time $\Theta(1)$.
- Bad: Space inefficient. $10^8 = 100 \text{ million slots allocated to store} \approx 500 \text{ K records}$.

Linked List:

- Good: Highly space efficient.
- Bad: *Search* runs slowly.

3.1.3 Hash Table Objectives

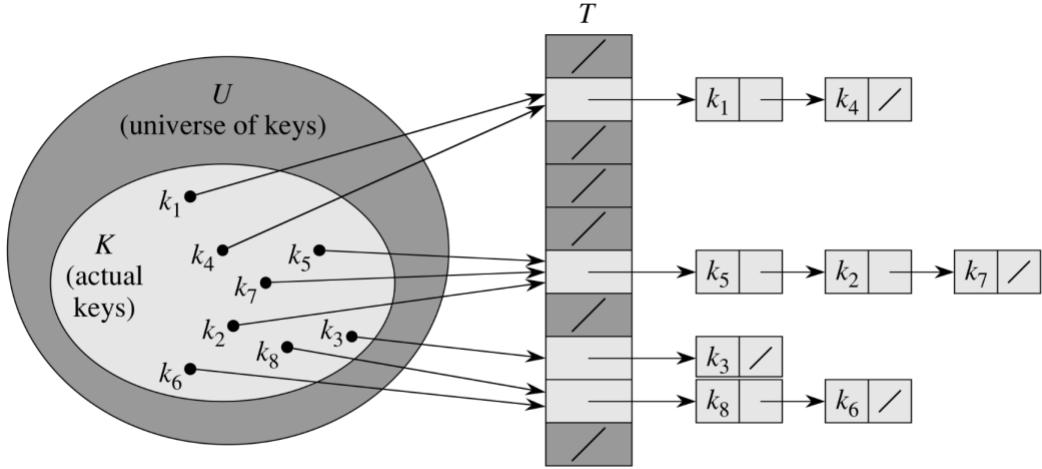
- Time Efficiency: *Search* to run in time $\approx \Theta(1)$.
- Space Efficiency: If we seek to store n records, space we should have to allocate $\approx \Theta(n)$.

3.1.4 Problem Setup

- We have a universe of keys, U .
- We seek to store n records each with a unique key. Typically, $n \ll |U|$.
- Designer of hash table knows: (i) U , and (ii) n , to some approximation.
- Designer of hash table does not know how the n keys are chosen from U .

3.1.5 How a Hash Table Works

- Designer picks some space allocation, m . Typically, $m = \Theta(n)$. Each $\{0, 1, \dots, m - 1\}$ is called a bucket.
- Designer picks a *hash function*, $h: U \rightarrow \{0, \dots, m - 1\}$.
- Designer picks a strategy for *resolving collisions*.
 - Two distinct keys, $k, l \in U$ are said to collide under h if $h(k) = h(l)$.
 - Collision guaranteed if $|U| > m$.
 - Can be much worse: if $|U| \geq nm$, there exists a set of keys $K \subseteq U$ with $|K| = n$ that all hash to the same bucket, immaterial of choice of h .
 - We consider collision resolution by *chaining* only.
 - If two records with keys k_i, k_j need to exist in the hash table, and $h(k_i) = h(k_j)$, then we maintain a linked list of the records in the bucket $h(k_i)$.



Credit: Cormen et al., "Introduction to Algorithms."

CREATENEW(m)

Pick a hash function, $h: U \rightarrow [0, m)$
 Create a table $T[0, \dots, m - 1]$
 Each $T[i] \leftarrow \text{NEWLIST}()$
 Store h as part of T
return T

DELETE(T, k)

Linked list $L \leftarrow T[h(k)]$
DELETEFROMLIST(L, k)

Choice of h is critical to performance. What is an “ideal” h ?

- Can be computed in time constant in n .
- Minimizes collisions. More precisely, acts like a “random function.”
 - Suppose for all distinct $k, l \in U$, $\Pr\{h(k) = h(l)\} = 1/m$.
 - Then in expectation, *length of each chain* = n/m .
 - Then, if $m = \Theta(n)$, expected time for *Search* is $\Theta(1)$.
 - Note: Worst case time for *Search* remains $\Theta(n)$.

INSERT(T, r)

Linked list $L \leftarrow T[h(key(r))]$
INSERTINTOLIST(L, r)

SEARCH(T, k)

Linked list $L \leftarrow T[h(k)]$
return **SEARCHINLIST(L, k)**

3.1.6 Universal Hashing

Given U and m , construct a set \mathcal{H} of hash functions with the following property:

Given any two distinct $k, l \in U$, the number of hash functions in \mathcal{H} for which k and l hash to the same bucket is $\leq |\mathcal{H}|/m$.

Then, in *CreateNew()*, pick an h from \mathcal{H} uniformly at random.

Claim 1: Given a universal set of hash functions, \mathcal{H} , if we pick $h \in \mathcal{H}$ uniformly at random, then for all distinct $k, l \in U$, $\Pr\{h(k) = h(l)\} \leq 1/m$.

Proof.

$$\Pr\{h(k) = h(l)\} = \frac{\text{\# hash functions in } \mathcal{H} \text{ for which } k \text{ and } l \text{ collide}}{\text{total \# hash functions in } \mathcal{H}}$$

$$\Pr\{h(k) = h(l)\} \leq \frac{|\mathcal{H}|/m}{|\mathcal{H}|} = 1/m$$

Challenge: How to construct such an \mathcal{H} given U, m .

3.1.7 How to Construct \mathcal{H} Given U, m

We use Student ID as keys: $U = \text{set of 8 digit numbers}$.

Now construct \mathcal{H} as follows:

- Generate some prime p such that $U \subseteq \{0, 1, \dots, p - 1\}$.
- For $a \in \{1, 2, \dots, p - 1\}, b \in \{0, 1, \dots, p - 1\}$, let $h_{a,b}(k) = ((ak + b) \bmod p) \bmod m$.
- Adopt $\mathcal{H} = \{h_{a,b} \mid a \in \{1, 2, \dots, p - 1\}, b \in \{0, 1, \dots, p - 1\}\}$.

Observation right off the bat: $|\mathcal{H}| = p(p - 1)$.

For the following claims, adopt:

- $k, l \in \{0, 1, \dots, p - 1\}, k \neq l$.
- a, b chosen uniformly (at random) from their respective sets.
- $r = (ak + b) \bmod p, s = (al + b) \bmod p$.

Claim 2: $k \neq l \rightarrow r \neq s$.

Proof. $r - s \equiv a(k - l) \pmod{p}$, and neither a nor $k - l \equiv 0 \pmod{p}$. And so, $r - s \not\equiv 0 \pmod{p} \neq 0$.

Claim 3: Let:

$$a, c \in \{1, \dots, p - 1\}$$

$$b, d \in \{0, \dots, p - 1\}$$

$$r_{a,b} = (ak + b) \bmod p$$

$$s_{a,b} = (al + b) \bmod p$$

$$r_{c,d} = (ck + d) \text{ mod } p$$

$$s_{c,d} = (cl + d) \text{ mod } p$$

Then: $\langle a, b \rangle \neq \langle c, d \rangle \Rightarrow \langle r_{a,b}, s_{a,b} \rangle \neq \langle r_{c,d}, s_{c,d} \rangle$

Proof. *

$$a = ((r_{a,b} - s_{a,b})((k-l)^{-1} \text{ mod } p)) \text{ mod } p$$

$$b = (r_{a,b} - ak) \text{ mod } p$$

and similarly for c and d

Therefore, $\langle r_{a,b}, s_{a,b} \rangle = \langle r_{c,d}, s_{c,d} \rangle \Rightarrow \langle a, b \rangle = \langle c, d \rangle$, a contradiction.

*Example: $8a + b = 5$ and $3a + b = 13$

These are both linear equations with unknowns a, b . Now, k, l are between 1 and $p - 1$, and since p is prime, $k - l$ is relatively prime to p , and so its multiplicative inverse exists.

Thus, for distinct $k, l \in U$:

- Each distinct $\langle a, b \rangle$ maps to distinct $\langle r, s \rangle$
- In each such $\langle r, s \rangle$, $r \neq s$.

Now:

- # distinct $\langle a, b \rangle$: $p(p - 1)$.
- # distinct r, s with $r \neq s$: $p(p - 1)$.
- So $\langle a, b \rangle$ to $\langle r, s \rangle$ mapping is one-to-one.
- So, picking $\langle a, b \rangle$ uniformly is equivalent to picking $\langle r, s \rangle$ uniformly.

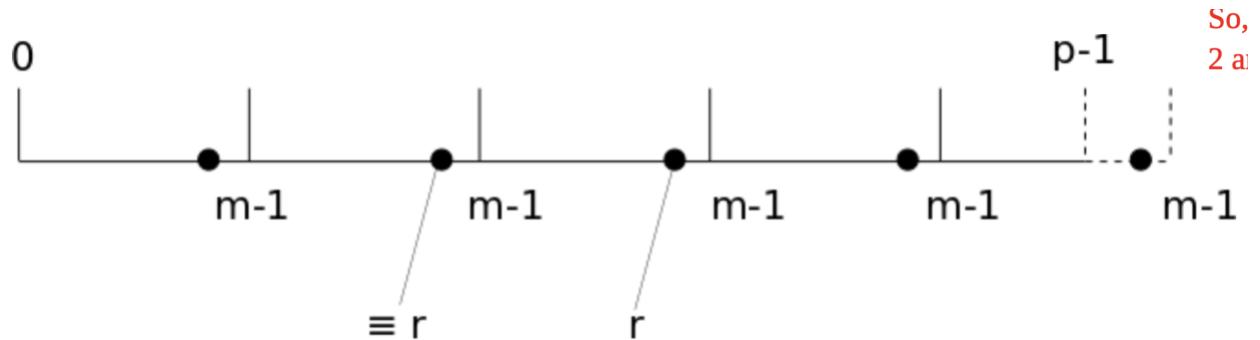
Finally, the $\text{mod } m$ part:

- Suppose each of a, b is chosen uniformly at random.
- Then, for $k \neq l$, $\Pr\{h_{a,b}(k) = h_{a,b}(l)\} = \Pr\{r \equiv s \pmod{m}\}$, r, s chosen uniformly.

$$\Pr\{r \equiv s \pmod{m}\} = \frac{\# \text{ items } \equiv r \pmod{m} \text{ but } \neq r \text{ in } \{0, \dots, p-1\}}{\# \text{ items } \neq r \text{ in } \{0, \dots, p-1\}}$$

$$\Pr\{r \equiv s \pmod{m}\} \leq \frac{\lfloor p/m \rfloor - 1}{p-1}$$

To intuit the $\lfloor p/m \rfloor$ term, pictorially:



*In the figure, $[p/m] = 5$. The value r , and every member of $\{0, 1, \dots, p - 1\} \neq r$, but $\equiv r \pmod{m}$ is shown as a bold dot.

And we know that $[x/y] \leq (x - 1)/y + 1$. So:

$$\Pr\{r \equiv s \pmod{m}\} \leq \frac{((p-1)/m + 1) - 1}{p-1} = 1/m$$

*Example: $p = 7, m = 3, r = 5$.

Then, $r \pmod{m} = 5 \pmod{3} = 2$.

So, in the set $\{0, \dots, 6\}$, the two numbers, 2 and 5, are congruent modulo 3: $2 \equiv 5 \pmod{3}$.

3.2 Divide and Conquer

3.2.1 Definition

Divide and Conquer: An algorithm design strategy to which some problems lend themselves.

The *divide and conquer* strategy solves a problem by:

1. Breaking it into *subproblems* that are themselves smaller instances of the same type of problem
2. Recursively solving these subproblems
3. Appropriately combining their answers

Example: Compute the product xy given $x, y \in \mathbb{Z}^+$.

Assume our encoding is binary. First split the bits of each x, y into the leading half the bits and trailing half the bits. E.g., $100011 = 100 \circ 011$.*

$$\begin{aligned} x &= x_L \circ x_R = 2^{n/2}x_L + x_R \\ y &= y_L \circ y_R = 2^{n/2}y_L + y_R \end{aligned}$$

Then:

$$\begin{aligned} xy &= 2^n x_L y_L + 2^{n/2} (x_L y_R + x_R y_L) + x_R y_R \\ xy &= 2^n x_L y_L + 2^{n/2} ((x_L + x_R)(y_L + y_R) - x_L y_L - x_R y_R) + x_R y_R \end{aligned}$$

*Example: In base 10, e.g., $153803 = 153 \times 10^{6/2} + 803$

So, to compute xy , where each x, y is n bits:

- Three multiplications of numbers $n/2$ bits long
- Two left bit shifts, one of n bits and the other of $n/2$ bits
- Six additions/subtractions each of n bits

Recurrence for running time, $T(n)$:

$$T(n) = \begin{cases} O(1), & \text{if } n = 1 \\ 3T(n/2) + O(n), & \text{otherwise} \end{cases}$$

To solve the recurrence, draw a recurrence tree, which is equivalent to inductive “string rewriting” as follows.

First, for the $O(n)$ term, we adopt a canonical function, n , and for $O(1)$, we adopt 1.

$$\begin{aligned} T(n) &= 3T\left(\frac{n}{2}\right) + n \\ &= 3\left(3T\left(\frac{n}{2^2}\right) + \frac{n}{2}\right) + n = 3^2 \cdot T\left(\frac{n}{2^2}\right) + n + \left(\frac{3}{2}\right)n \\ &= 3^3 \cdot T\left(\frac{n}{2^3}\right) + n \cdot \left(\left(\frac{3}{2}\right)^0 + \left(\frac{3}{2}\right)^1 + \left(\frac{3}{2}\right)^2\right) \\ &= 3^4 \cdot T\left(\frac{n}{2^4}\right) + n \cdot \left(\left(\frac{3}{2}\right)^0 + \left(\frac{3}{2}\right)^1 + \left(\frac{3}{2}\right)^2 + \left(\frac{3}{2}\right)^3\right) \\ &\quad \dots \\ &= 3^{\log_2 n} \cdot 1 + n \cdot \left(\left(\frac{3}{2}\right)^0 + \left(\frac{3}{2}\right)^1 + \dots + \left(\frac{3}{2}\right)^{\log_2 n-1}\right) (1, 2, 3) \end{aligned}$$

$$\begin{aligned}
&= n^{\log_2 3} + 2n \left(\left(\frac{3}{2}\right)^{\log_2 n} - 1 \right) = 3 \cdot n^{\log_2 3} - \frac{2}{n} \\
&= O(n^{1.59})
\end{aligned}$$

(1): For what x is $\frac{n}{2^x} = 1$? Answer: $\frac{n}{2^x} = 1$ iff $n = 2^x$ iff $\log_2 n = x$

(2): $3^{\log_2 n} = z$, $\log_2 z = \log_2 n \times \log_2 3$, then exponentiate by 2: $z = (2^{\log_2 n})^{\log_2 3} = n^{\log_2 3}$

(3): $S = a^0 + a^1 + a^2 + \dots + a^y \rightarrow aS = a^1 + \dots + a^y + a^{y+1}$

Subtract: $(1 - a)S = a^0 - a^{y+1}$

Implies: $S = \frac{a^0 - a^{y+1}}{1-a}$

3.2.2 Solving Recurrences

Suppose we have the following recurrence for a function $f: \mathbb{N} \rightarrow \mathbb{R}^+$, with constants $a, b, d \in \mathbb{R}, a > 0, b > 1, d \geq 0$:

$$f(n) = \begin{cases} \text{constant, if } n \text{ is small} \\ af(\lceil n/b \rceil) + O(n^d), \text{ otherwise} \end{cases}$$

Claim 1 (Master Theorem): Given $f(n)$ as in the recurrence above,

$$f(n) = \begin{cases} O(n^d), \text{ if } d > \log_b a \\ O(n^d \log n), \text{ if } d = \log_b a \\ O(n^{\log_b a}), \text{ otherwise, i.e., } d < \log_b a \end{cases}$$

Proof. Draw a tree and count.

Example 1: Binary search $\text{BinSearch}(A[1, \dots, n], i)$ on array of n items. Running time, $T(n)$, as measured by # comparisons of items in the (sorted) input array and the item we're searching for:

$$T(n) = \begin{cases} 1, \text{ if } n = 1 \\ T(n/2) + 1, \text{ otherwise} \end{cases}$$

We have $a = 1, b = 2, d = 0$. So, $d = \log_b a$, and $T(n) = O(\log n)$.

Example 2: Recurrence for our algorithm for multiplication.

$$T(n) = \begin{cases} O(1), \text{ if } n = 1 \\ 3T(n/2) + O(n), \text{ otherwise} \end{cases}$$

We have $a = 3, b = 2, d = 1$. So, $d < \log_b a$, and $T(n) = O(n^{\log_2 3})$.

3.2.3 Sorting Problem

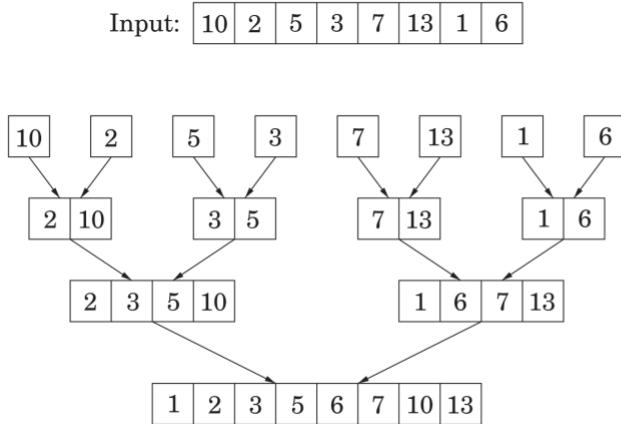
Given as input an array $A[1, \dots, n]$ whose items are drawn from a totally ordered set, return the sorted permutation of A .

```
function mergesort(a[1...n])
Input: An array of numbers a[1...n]
Output: A sorted version of this array

if n > 1:
    return merge(mergesort(a[1...[n/2]]),
    mergesort(a[[n/2] + 1...n]))
else:
    return a
```

```
function merge(x[1...k], y[1...l])
if k = 0: return y[1...l] We claim that a count of this comparison
if l = 0: return x[1...k] is a meaningful measure of running-time
if x[1] ≤ y[1]: _____ of mergesort
    return x[1] ∘ merge(x[2...k], y[1...l])
else:
    return y[1] ∘ merge(x[1...k], y[2...l])
```

Figure 2.4 The sequence of merge operations in `mergesort`.



Claim 2: *merge* is correct.

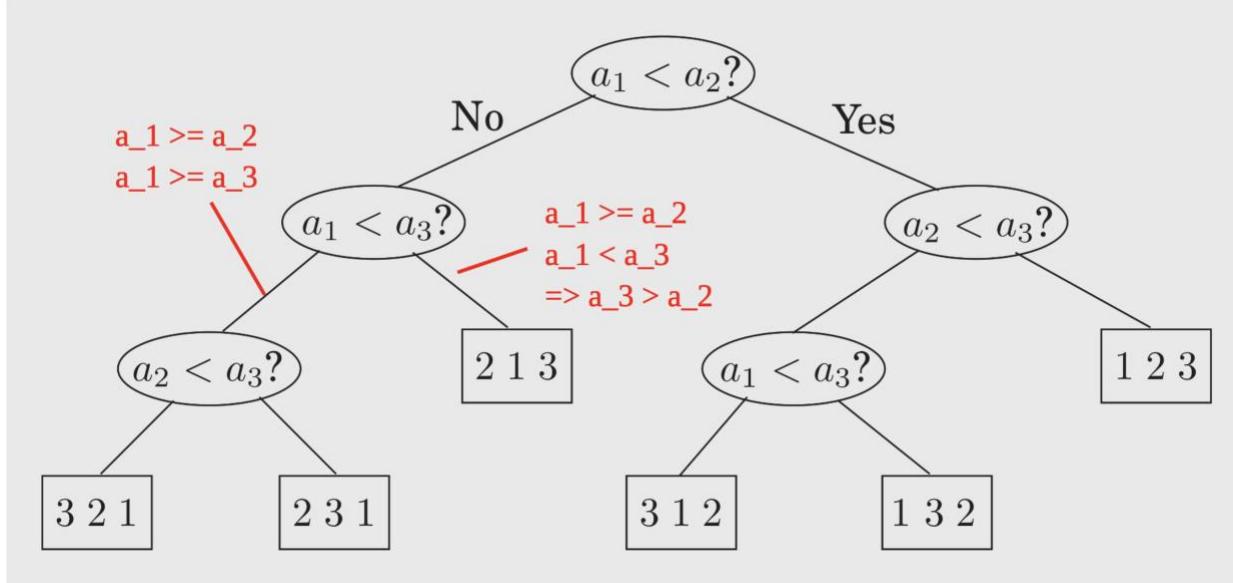
Proof. Induction on # recursive calls.

Claim 3: *mergesort* is correct.

Proof. Induction on # recursive calls. Appeal to Claim 2.

$$\# \text{ comparisons of array items: } T(n) = 2T\left(\frac{n}{2}\right) + O(n) \Rightarrow T(n) = O(n \log n).$$

Claim 4: *mergesort* is optimal in the # comparisons: any comparison-based sorting algorithm performs $\Omega(n \log n)$ comparisons.



- What is the most compact way possible to consider every possible permutation? Where “way possible” refers to pairwise comparisons.
- Answer: Binary tree as shown above, with every permutation showing up as a leaf.
- Minimum # comparisons (in worst case) = longest path from root to a leaf = $\log(n!) = \Omega(n \log n)$.
- One of the rare instances in which we have a lower bound for an algorithmic problem.

3.3 Divide and Conquer Applications

3.3.1 Selection Problem

Given as input: (i) an array $S[1, \dots, n]$ of items drawn from a totally ordered set, and (ii) an integer $k \in \{1, \dots, n\}$, identify the k^{th} smallest item in S .

E.g., if $k = 1$, we seek the smallest. If $k = n$, we seek the largest. If $k = \lfloor (n + 1)/2 \rfloor$, we seek the (left) median.

A divide and conquer strategy:

- Pick some member of S as the *pivot*, v .
 - E.g., randomly from amongst all items in S .
- *Partition* or *split* around that pivot v .
 - Rearrange items in S as follows:

- Move every item $< v$ to the left of v .
- Move every item $> v$ to the right of v .
- Thus, every item whose value is v ends up where it would in a sorted permutation of S .
- Check and recurse.

$$S : \boxed{2 \ 36 \ 5 \ 21 \ 8 \ 13 \ 11 \ 20 \ 5 \ 4 \ 1}$$

is split on $v = 5$, the three subarrays generated are

$$S_L : \boxed{2 \ 4 \ 1} \quad S_v : \boxed{5 \ 5} \quad S_R : \boxed{36 \ 21 \ 8 \ 13 \ 11 \ 20}$$

$$\text{selection}(S, k) = \begin{cases} \text{selection}(S_L, k) & \text{if } k \leq |S_L| \\ v & \text{if } |S_L| < k \leq |S_L| + |S_v| \\ \text{selection}(S_R, k - |S_L| - |S_v|) & \text{if } k > |S_L| + |S_v|. \end{cases}$$

If v chosen to be, e.g., median, then $T(n) = T(n/2) + O(n) \Rightarrow T(n) = O(n)$.

Challenge is now: How to guarantee a good choice for v ?

Suppose “good” v is some value in the middle 50% of S .

- E.g., if $n = 30$, $8 \leq \text{rank}[v] \leq 22$.

Then, recurrence for running time of *selection*:

$$T(n) \leq T\left(\frac{3}{4} \cdot n\right) + O(n) = O(n)$$

To pick such a “good” v :

- Pick random v .
- Check if it is good: *time* = $O(n)$.
- If not good, repeat.

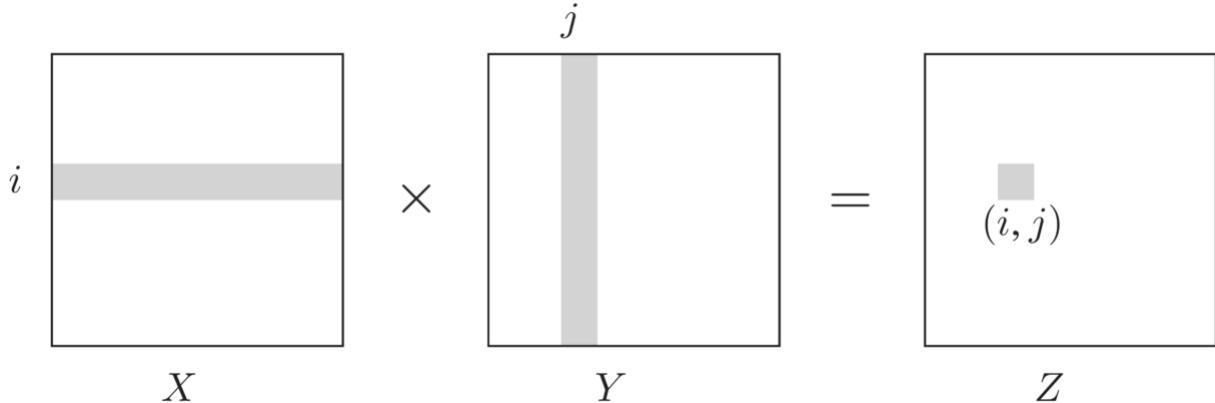
$$\Pr\{\text{we pick "good" } v \text{ in a trial}\} = \frac{1}{2}$$

So # trials in expectation to get a “good” $v = 2$,

So, we have an algorithm for *selection*, which runs in time $O(n)$ in the expected case.

3.3.2 Matrix Multiplication

Given two $n \times n$ matrices, X, Y , compute the $n \times n$ product $Z = X \cdot Y$.



$$Z = [z_{i,j}] \text{ where } z_{i,j} = \sum_{k=1}^n x_{i,k} \cdot y_{k,j}$$

Naïve algorithm to compute Z : $\Theta(n^3)$ scalar multiplications and additions.

Smarter algorithm based on divide-n-conquer:

each of A, B, C, D is $n/2 \times n/2$

$$X = \begin{bmatrix} A & B \\ C & D \end{bmatrix}, \quad Y = \begin{bmatrix} E & F \\ G & H \end{bmatrix}$$

$$XY = \begin{bmatrix} AE + BG & AF + BH \\ CE + DG & CF + DH \end{bmatrix}$$

$$= \begin{bmatrix} P_5 + P_4 - P_2 + P_6 & P_1 + P_2 \\ P_3 + P_4 & P_1 + P_5 - P_3 - P_7 \end{bmatrix}$$

where

$$\begin{aligned} P_1 &= A(F - H) \\ P_2 &= (A + B)H \\ P_3 &= (C + D)E \\ P_4 &= D(G - E) \end{aligned}$$

$$\begin{aligned} P_5 &= (A + D)(E + H) \\ P_6 &= (B - D)(G + H) \\ P_7 &= (A - C)(E + F) \end{aligned}$$

Now, recurrence for # scalar multiplications and additions:

$$\begin{aligned} T(n) &= 7T\left(\frac{n}{2}\right) + O(n^2) \\ &= O(n^{\log_2 7}) \approx O(n^{2.81}) \end{aligned}$$

4 DECOMPOSITIONS OF GRAPHS, PATHS

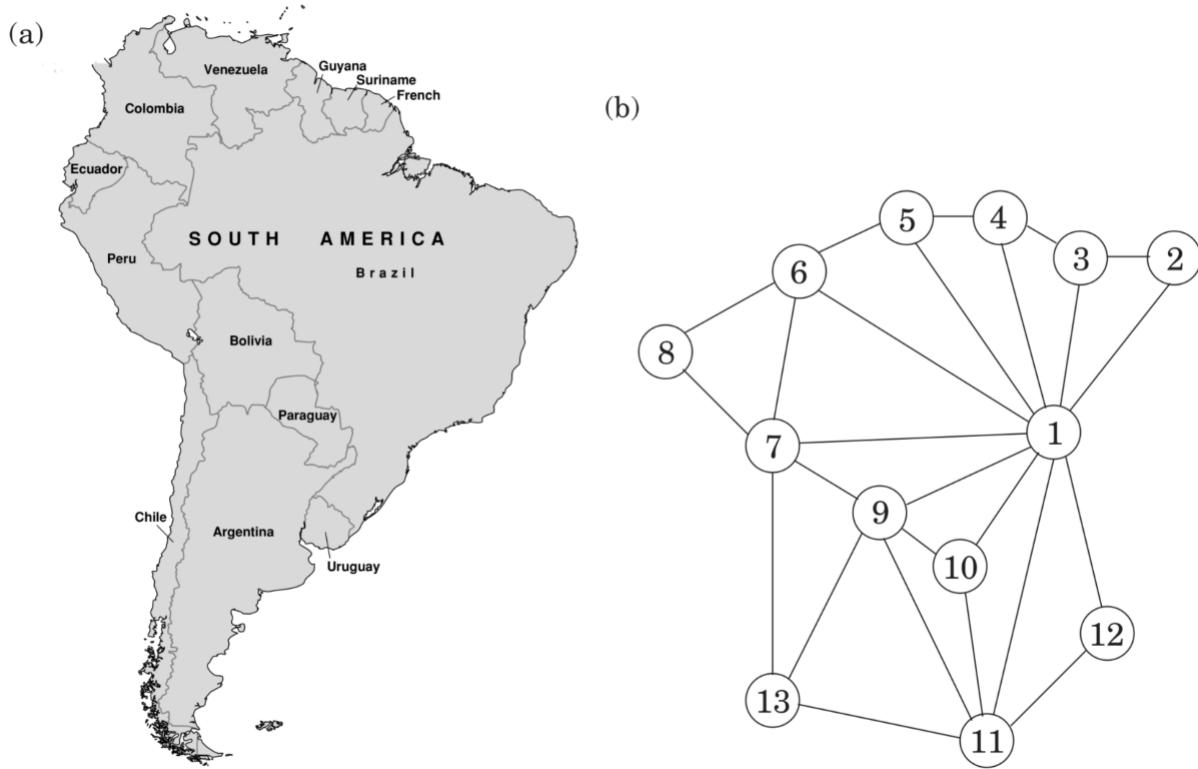
4.1 Graphs

Graph $G = \langle V, E \rangle$. $V = \text{set of vertices}$. $E = \text{set of edges}$.

- Directed: $E \subseteq V \times V$. Written $\langle u, v \rangle \in E$.
- Undirected: $E \subseteq \{\{u, v\} \mid u, v \in V, u \neq v\}$. Abuse notation and write $\langle u, v \rangle \in E$.

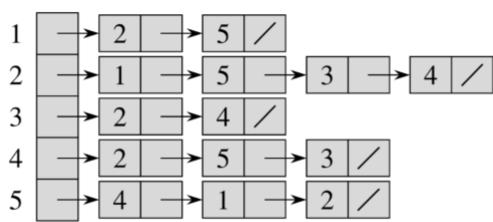
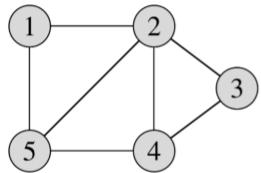
Lots of things are modelled using graphs. E.g., maps, digital circuits, social networks.

Figure 3.1 (a) A map and (b) its graph.

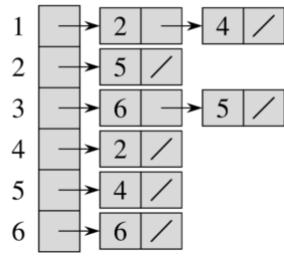
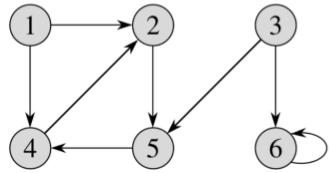


4.1.1 Representing a Graph

Adjacency list and adjacency matrix.



	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0



	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

4.1.2 Graph Exploration

An encoding of a graph can be seen as local information, i.e., just information about adjacency.

Algorithmic problems with graphs ask about some global property of the graph.

Given a graph $G = \langle V, E \rangle$ and two distinct vertices $u, v \in V$, is v reachable from u ?

An algorithm for this: Depth-First Search (DFS).

Figure 3.5 Depth-first search.

```
procedure dfs(G)
for all  $v \in V$ :
    visited( $v$ ) = false
for all  $v \in V$ :
    if not visited( $v$ ): explore( $v$ )
```

Figure 3.3 Finding all nodes reachable from a particular node.

```
procedure explore( $G, v$ )
Input:  $G = \langle V, E \rangle$  is a graph;  $v \in V$ 
Output: visited( $u$ ) is set to true for all nodes  $u$  reachable
       from  $v$ 

visited( $v$ ) = true
previsit( $v$ )
for each edge  $(v, u) \in E$ :
    if not visited( $u$ ): explore( $u$ )
postvisit( $v$ )
```

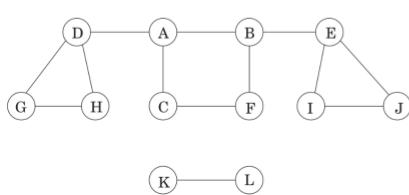
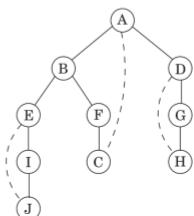


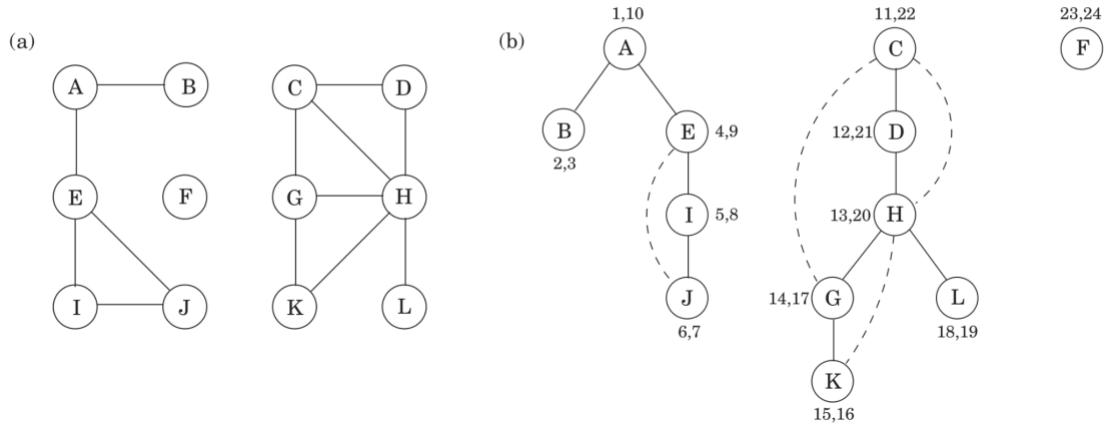
Figure 3.4 The result of explore(A) on the graph of Figure 3.2.



Claim 1: v is reachable from y if and only if $\text{explore}(G, u)$ causes v to be **visited**.

Time efficiency: $O(|V| + |E|)$.

Figure 3.6 (a) A 12-node graph. (b) DFS search forest.



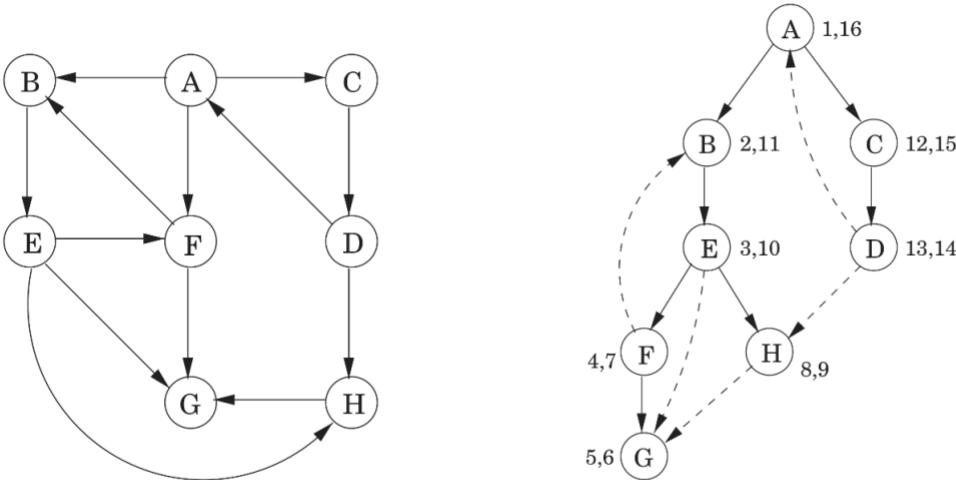
```
procedure previsit( $v$ )
ccnum[ $v$ ] = cc
```

```
procedure previsit( $v$ )
pre[ $v$ ] = clock
clock = clock + 1
```

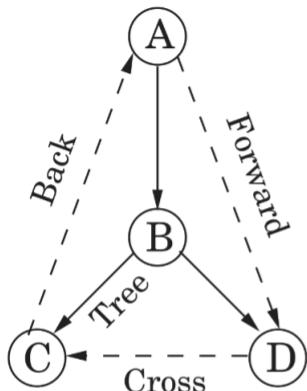
```
procedure postvisit( $v$ )
post[ $v$ ] = clock
clock = clock + 1
```

4.2 DFS, Continued

Figure 3.7 DFS on a directed graph.

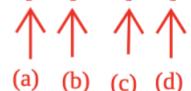


DFS tree



pre/post ordering for (u, v)	Edge type
$[u \quad v \quad v \quad u]$	Tree/forward
$[v \quad u \quad u \quad v]$	Back
$[v \quad v \quad u \quad u]$	Cross

Specifically: if $\langle u, v \rangle$ is an edge, then $[u \quad v \quad u \quad v]$ cannot happen.



(a): explore(u) pushed onto call stack

(b): explore(v) pushed, note: explore(u) is lower in the stack, unpopped

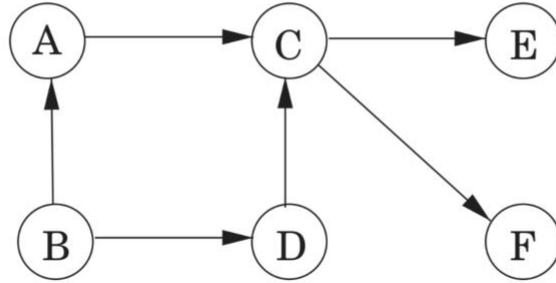
(c): explore(u) popped off call stack

(d): explore(v) popped off call stack

4.2.1 Directed Acyclic Graph (DAG)

Linearization or topological sort: of a DAG is a left-to-right ordering of all the vertices, $\langle u_0, u_1, \dots, u_{|V|-1} \rangle$ such that $\langle u_i, u_j \rangle \in E \Rightarrow j > i$, i.e., all edges go left to right.

Figure 3.8 A directed acyclic graph with one source, two sinks, and four possible linearizations.



A linearization: B, D, A, C, F, E .

Claim 1: If G is a DAG, then in a run of $\text{dfs}(G)$, every edge leads to a vertex with a lower (i.e., smaller-valued) **post** number.

Equivalent to: if G is a DAG, then in any run of $\text{dfs}(G)$, no edge is a back edge.

An algorithm for linearization:

Run $\text{dfs}(G)$

When we are done with a vertex, insert it at the beginning.

Claim 2: Directed G is acyclic if and only if it is linearizable if and only if $\text{dfs}(G)$ results in no back edges.

Another algorithm:

Find a source, output it, and delete it from the graph.

Repeat until the graph is empty.

Claim: This algorithm runs in time $O(|V| + |E|)$ as well.

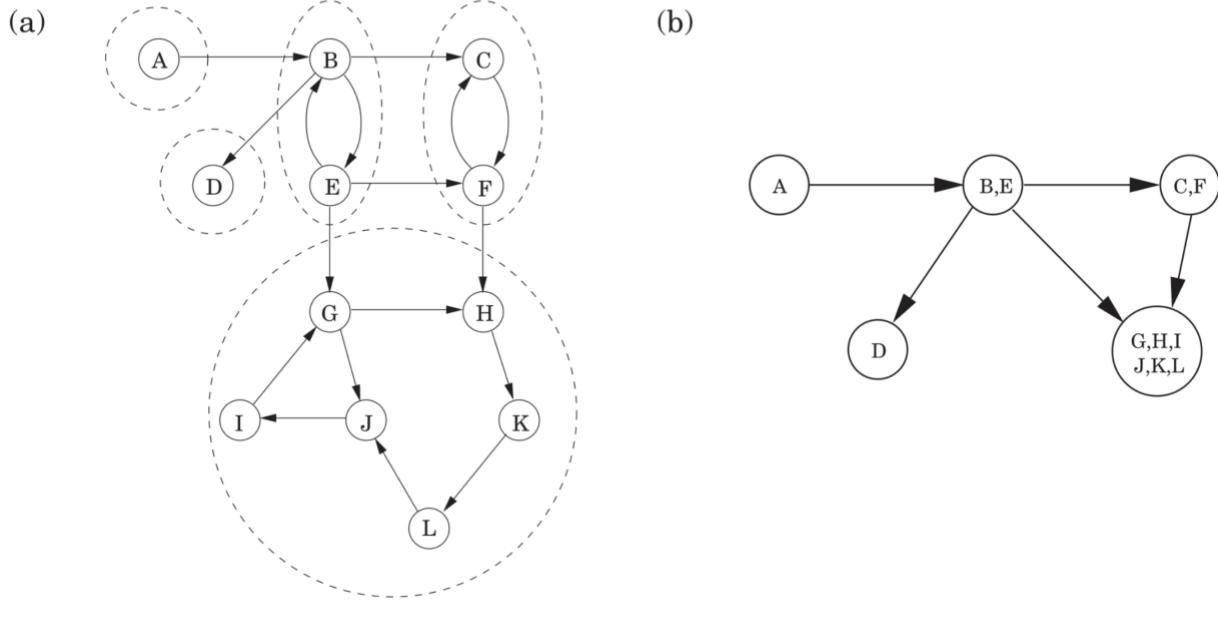
4.2.2 Strongly Connected Components

Given directed $G, u, v \in G$ with $u \neq v$ are said to be strongly connected if $u \rightsquigarrow v$ and $v \rightsquigarrow u$.

Given directed $G = \langle V, E \rangle$, the set of Strongly Connected Components $\{V_0, V_1, \dots, V_{k-1}\}$ is a partition of V with the property: for all $i = 0, \dots, k - 1$, distinct $u, v \in V_i \Leftrightarrow u \rightsquigarrow v$ and $v \rightsquigarrow u$.

Recall: A partition of a set $S, \{S_0, \dots, S_{n-1}\}$ has the properties: (i) for all $i = 0, \dots, n - 1, S_i \subseteq S$, and (ii) for all distinct $i, j = 0, \dots, n - 1, S_i \cap S_j = \emptyset$, and (iii) $S_0 \cup S_1 \cup \dots \cup S_{n-1} = S$.

Figure 3.9 (a) A directed graph and its strongly connected components. (b) The meta-graph.

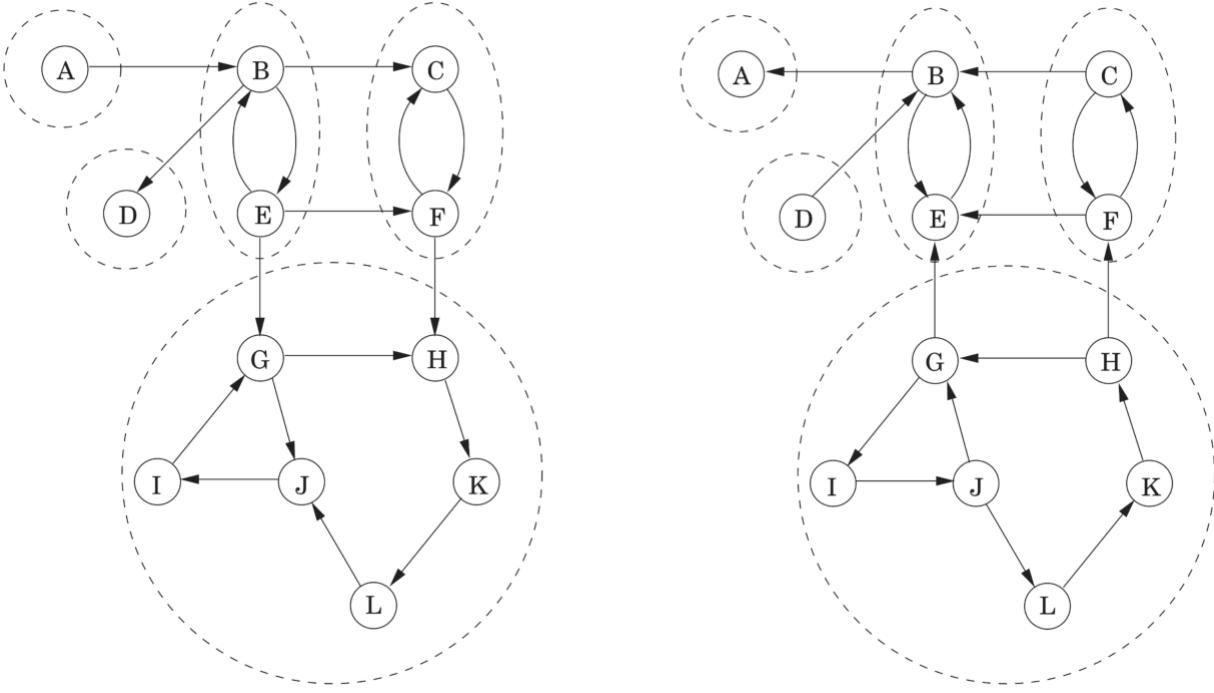


Claim 3: The meta-graph is (i) unique, and (ii) acyclic.

The algorithm for computing the meta-graph:

1. Run **dfs** on G^R .
2. Run the undirected connected components algorithms on G , and during the depth-first search, process the vertices in decreasing order of their **post** numbers from Step (1).

Where G^R is the transpose or reverse graph of G .



Claim 4: G^R can be computed from $G = \langle V, E \rangle$ in time $O(|V| + |E|)$.

Claim 5: C is a SCC in G if and only if C is a SCC in G^R .

Definition 1: In a directed graph $G = \langle V, E \rangle$, $u \in V$ is called a *source* if there is no edge that is incident on u . A node $v \in V$ is called a *sink* if there is no edge that leaves v .

Claim 6: A node is a source in G if and only if it is a sink in G^R . A node is a sink in G if and only if it is a source in G^R .

Proof. Exploit the sub-claim: $(G^R)^R = G$.

Claim 7: A node that receives the highest *post* number in a run of **dfs** on a directed graph F lies in a source in SCC of the meta-graph.

- And therefore, a sink SCC of F^R .

Claim 8: If C, C' are distinct SCCs in F , and there's an edge $u \rightarrow v$ where $u \in C, v \in C'$, then the highest *post* number in C is bigger than the highest *post* number in C' in any run of **dfs** on F .

- A generalization of Claim (7) above.

Step (1) of the algorithm on G^R example above can result in: $G, I, J, L, K, H, D, C, F, B, E, A$.

Then, Step (2) may group, in sequence: $(G, H, I, J, K, L), (D), (C, F), (B, E), (A)$.

Sanity-check question: if that was the sequence of **post** numbers in Step (1), which vertex happened to have been chosen first in **dfs**? More generally, what was the order of vertices chosen in **dfs**?

- Solution: Work backwards from the last one, i.e., A .
- First vertex chosen was A . Then B . Then C . Then D . Finally, G .

4.3 Shortest Paths

- Another “global” property in the context of a graph.
- A path (of length k) in $G = \langle V, E \rangle$ is a sequence $\langle u_0, u_1, \dots, u_k \rangle$ where $\langle u_i, u_{i+1} \rangle \in E$ for all $i = 0, \dots, k - 1$. A path is simple if all vertices in it are distinct.
- A shortest path from vertex u to v is one of minimum length across all paths $u \rightsquigarrow v$. We call the length of the shortest path $u \rightsquigarrow v$ the distance from u to v .
- A shortest path $u \rightsquigarrow v$ is not necessarily unique. The length of a shortest path $u \rightsquigarrow v$ is.

Claim 1: If a path $u \rightsquigarrow v$ exists, then a shortest path $u \rightsquigarrow v$ exists. If a shortest path $u \rightsquigarrow v$ exists, it must be simple.

4.3.1 Breadth First Search (BFS)

- An algorithm to compute distances and shortest paths from a source vertex.
- Mindset: Think of edges as cloth strings of equal length. A BFS tree results from picking the graph up by the source vertex.

Figure 4.2 A physical model of a graph.

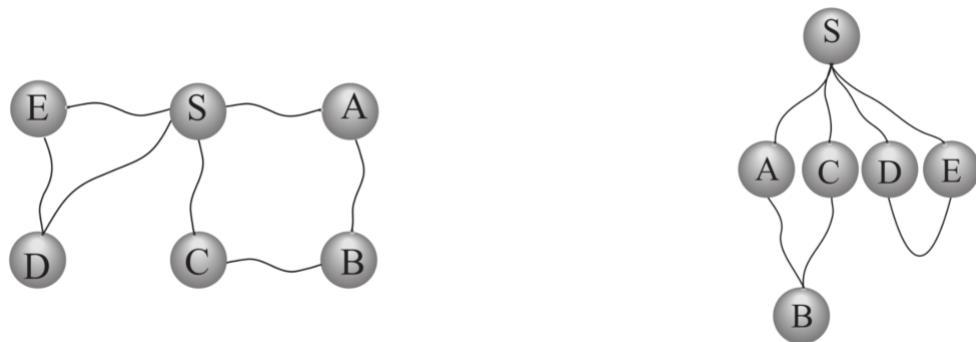


Figure 4.3 Breadth-first search.

```

procedure bfs( $G, s$ )
Input: Graph  $G = (V, E)$ , directed or undirected; vertex  $s \in V$ 
Output: For all vertices  $u$  reachable from  $s$ ,  $\text{dist}(u)$  is set
       to the distance from  $s$  to  $u$ .

for all  $u \in V$ :
     $\text{dist}(u) = \infty$ 

 $\text{dist}(s) = 0$ 
 $Q = [s]$  (queue containing just  $s$ )
while  $Q$  is not empty:
     $u = \text{eject}(Q)$ 
    for all edges  $(u, v) \in E$ :
        if  $\text{dist}(v) = \infty$ :
            inject( $Q, v$ )
             $\text{dist}(v) = \text{dist}(u) + 1$ 

```

Figure 4.1 (a) A simple graph and (b) its depth-first search tree.

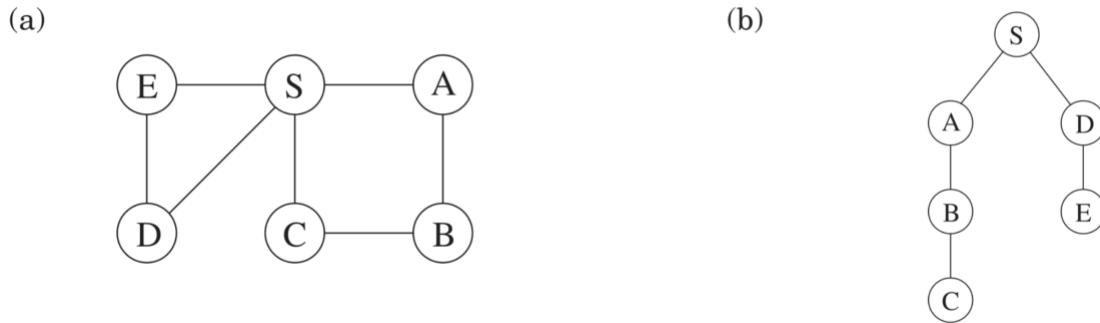
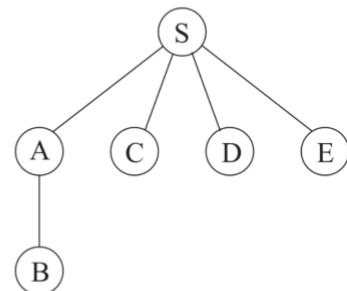


Figure 4.4 The result of breadth-first search on the graph of Figure 4.1.

Order of visitation	Queue contents after processing node
S	$[S]$
A	$[A C D E]$
C	$[C D E B]$
D	$[D E B]$
E	$[E B]$
B	$[B]$
	$[]$



Correctness of BFS:

Claim 2: For every $v \in V$, $\text{bfs}(G, s)$ results in the correct value for $\text{dist}(v)$.

Proof. Two cases.

- (i) No path $s \rightsquigarrow v$ exists. Then, for BFS to be correct for that particular v , $\text{dist}(v) = \infty$ at the end of the run of the algorithm.
- (ii) If a path $s \rightsquigarrow v$ exists, then let $d \in \mathbb{Z}_0^+$ be the distance from s to v . Then prove by induction on d that $\text{dist}(v) = d$. Use the fact that $\text{dist}(v)$, once initialized, is then set only when v is enqueued, and never again thereafter.

Time efficiency of BFS: $O(|V| + |E|)$.

Unlike the manner in which we designed DFS, not every vertex and edge is necessarily visited by BFS. So, its time efficiency is $\Theta(|V| + |E|)$ in the worst case only.

5 PATHS IN GRAPHS, CONTINUED

5.1 Shortest Paths in Weighted Graphs

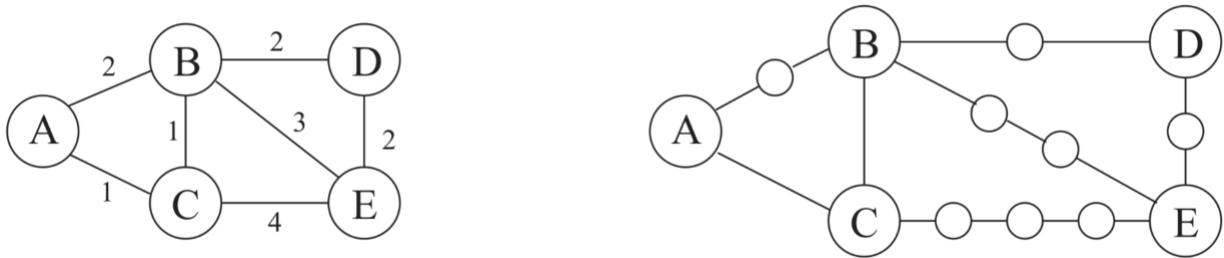
We now assume that a graph is a triple $G = \langle V, E, l \rangle$, where $l: E \rightarrow \mathbb{R}$.

Suppose $l_e \in \mathbb{Z}^+$ for all $e \in E$. Then, a candidate algorithm:

Introduce $l_e - 1$ “dummy” vertices on edge e .

Run bfs on the resultant graph.

Figure 4.6 Breaking edges into unit-length pieces.



Algorithm is correct, but inefficient. How inefficient?

Answer: In the worst case, algorithm runs in exponential time, i.e., in time exponential in the size of the input.

In binary encoding, each l_e can be encoded with $\Theta(\log l_e)$ bits only.

Better adaptation of BFS: Dijkstra's algorithm.

5.1.1 Dijkstra's Algorithm

An example of a greedy algorithm: Locally optimal choices result in global optimum.

Caution 1: Not every (optimization) problem possesses such a greedy choice property.

Caution 2: Not necessarily easy to identify such a greedy choice property even if problem possesses one.

Caution 3: Dijkstra's guaranteed to be correct if $l_e \in \mathbb{R}_0^+$ for all $e \in E$.

Dijkstra's assumes access to a priority queue data structure. API:

- *makequeue*: Make a priority queue with a given set of values. May be realized as repeated calls to an *insert* algorithm.

- *deletemin*: Remove a queue item of smallest value and return it.
- *decreasekey*: Decrease the value of an item in the queue.

Figure 4.8 Dijkstra's shortest-path algorithm.

```

procedure dijkstra( $G, l, s$ )
Input: Graph  $G = (V, E)$ , directed or undirected;
       positive edge lengths  $\{l_e : e \in E\}$ ; vertex  $s \in V$ 
Output: For all vertices  $u$  reachable from  $s$ ,  $\text{dist}(u)$  is set
       to the distance from  $s$  to  $u$ .

for all  $u \in V$ :
     $\text{dist}(u) = \infty$ 
     $\text{prev}(u) = \text{nil}$ 
 $\text{dist}(s) = 0$ 

 $H = \text{makequeue}(V)$  (using dist-values as keys)
while  $H$  is not empty:
     $u = \text{deletemin}(H)$ 
    for all edges  $(u, v) \in E$ :
        if  $\text{dist}(v) > \text{dist}(u) + l(u, v)$ :
             $\text{dist}(v) = \text{dist}(u) + l(u, v)$ 
             $\text{prev}(v) = u$ 
            decreasekey( $H, v$ )
```

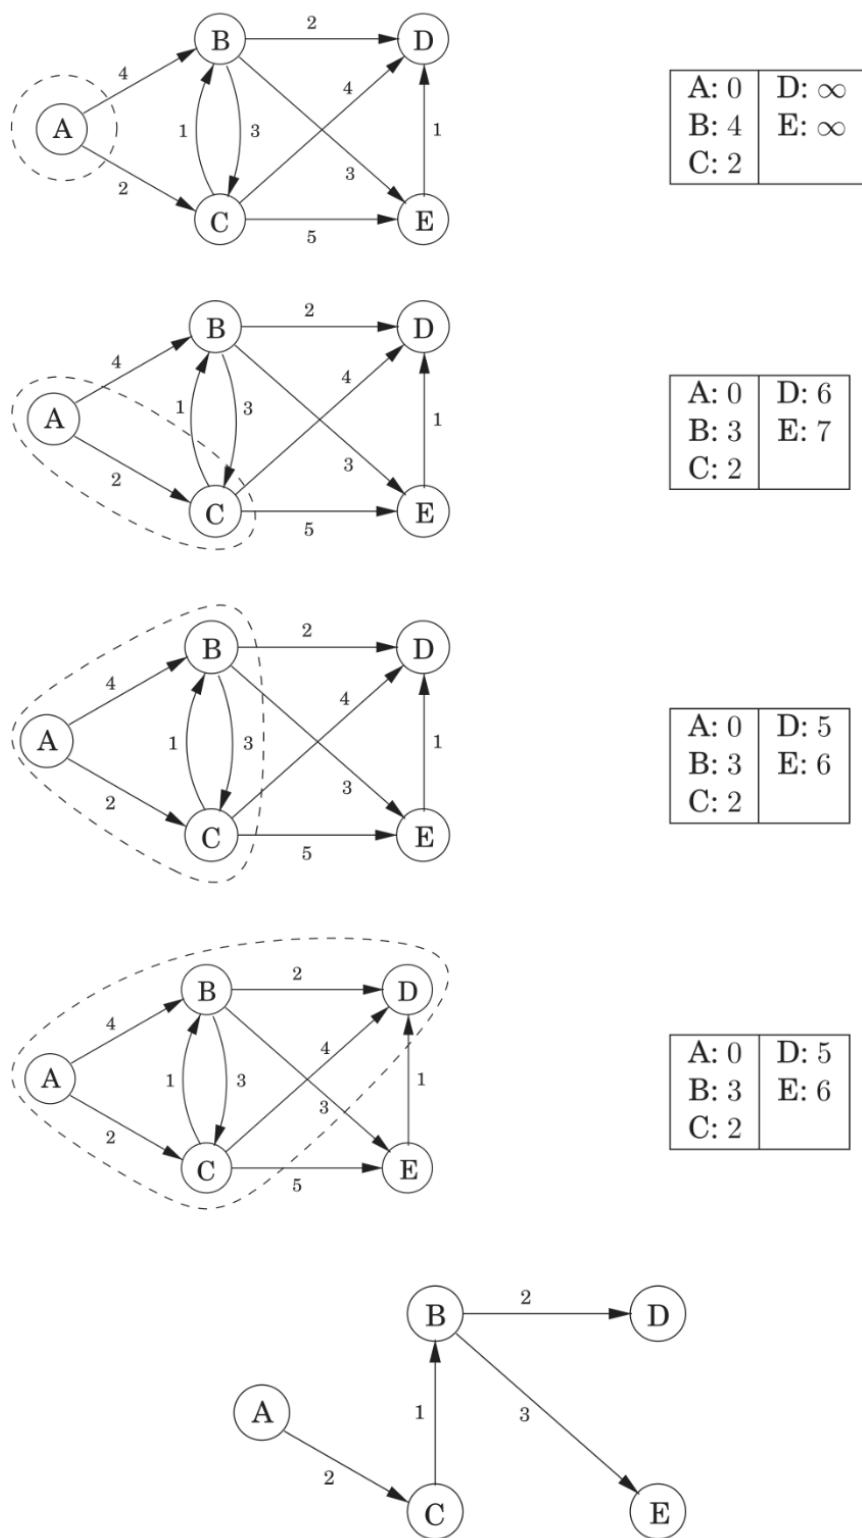
really, the dist values as the priority

"update" or "relaxation" of an edge $\langle u, v \rangle$

Semantics/properties of “ ∞ ”:

- $\infty = \infty$. Or $a = \infty$ and $b = \infty \Rightarrow a = b$.
- For any $r \in \mathbb{R}, r < \infty$.
- For any $r \in \mathbb{R}, r + \infty = \infty + r = \infty + \infty = \infty$.

Figure 4.9 A complete run of Dijkstra's algorithm, with node *A* as the starting point. Also shown are the associated *dist* values and the final shortest-path tree.



5.1.2 Correctness

A partial proof. First, denote as $\delta(x, y)$ the shortest distance from vertex x to y .

We prove:

For every vertex u that is reachable from s , at the moment u is extracted from the priority queue, it is the case that $\text{dist}(u) = \delta(s, u)$.

Claim 1: If $a \rightsquigarrow x \rightsquigarrow y \rightsquigarrow b$ is a shortest path from a to b , then the x to y subpath is a shortest path from x to y .

Proof. By contradiction.

Claim 2: Suppose $s \rightsquigarrow x \rightarrow y$ is a shortest path from s to y , where $x \rightarrow y$ is an edge. Also suppose that we are at a moment that $\text{dist}(x) = \delta(s, x)$. And we update $\text{dist}(y)$ based on $l(x, y)$ as in Dijkstra's algorithm. Then, immediately after, $\text{dist}(y) = \delta(s, y)$.

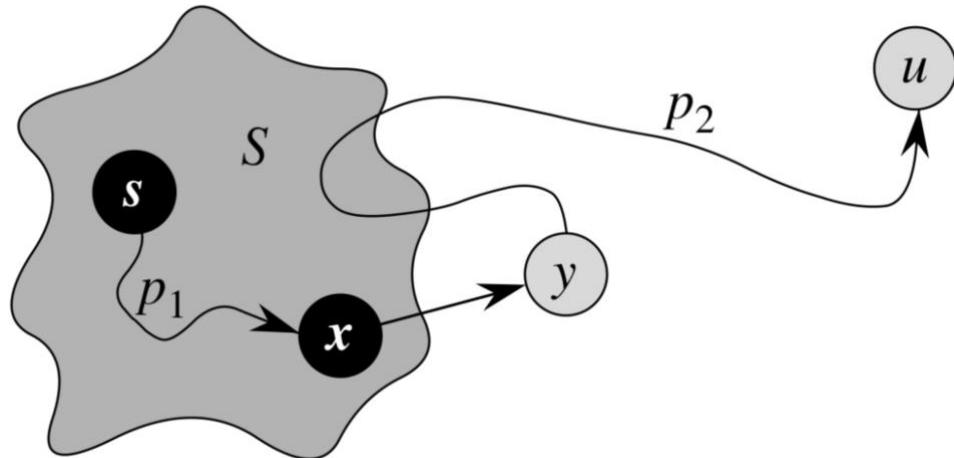
Proof. By Claim (1), each of the subpaths $s \rightsquigarrow x$ and $x \rightarrow y$ is a shortest path. Therefore, $\delta(s, y) = \delta(s, x) + l(x, y)$.

Claim 3: Once we initialize $\text{dist}(u)$ to ∞ , it is always the case that $\delta(s, u) \leq \text{dist}(u)$.

Note 1: Claims (1)-(3) hold even if we allow negative edge weights.

Note 2: Claim (1) refers to a kind of *optimal substructure*.

Now the proof for the claim about Dijkstra's:



Consider the moment in a run of Dijkstra's algorithm as shown above.

- The set S comprises those vertices that have already been extracted from the priority queue. At the minimum, $S = \{s\}$, but may contain additional vertices. We assume that for every vertex in S , its $dist$ value is indeed correct, i.e., equal to its $\delta(s, \cdot)$ value.
- We are about to extract the vertex u from the priority queue, i.e., its $dist$ value is the minimum in the priority queue at this moment. For the purpose of contradiction, assume that $dist(u) \neq \delta(s, u)$. That is, u is the first vertex in this run of Dijkstra's for which we have an error.
- Assume that the $s \rightsquigarrow u$ path shown in the picture is a shortest path from s to u . That path can be decomposed into: $s \xrightarrow{p_1} x \rightarrow y \xrightarrow{p_2} u$, where: (i) all the vertices of p_1 are $\in S$, (ii) y is the first vertex on the $s \rightsquigarrow u$ path that $\notin S$, (iii) we know $s \neq u$, and (iv) it is possible that $s = x$ and/or $y = u$.

We now derive a contradiction to the claim that $dist(u) \neq \delta(s, u)$.

- $dist(y) = \delta(s, y)$ – by Claim (2) above.
- $\delta(s, y) \leq \delta(s, u)$ – by Claim (1) above, and all edge weights are non-negative.
- $\delta(s, u) \leq dist(u)$ – by Claim (3) above.
- $dist(u) \leq dist(y)$ – we are about to extract u from the priority queue.

So, we have:

$$\begin{aligned} dist(y) &= \delta(s, y) \leq \delta(s, u) \leq dist(u) \leq dist(y)^* \\ \Rightarrow dist(y) &= \delta(s, y) = \delta(s, u) = dist(u) \end{aligned}$$

* $a \leq b \leq c \leq d \leq a \rightarrow a = b = c = d$

We have attained our desired contradiction.

5.1.3 Time Efficiency of Dijkstra

Implementation	<code>deletemin</code>	<code>insert/ decreasekey</code>	$ V \times \text{deletemin} +$ $(V + E) \times \text{insert or decreasekey}$
Array	$O(V)$	$O(1)$	$O(V ^2)$
Binary heap	$O(\log V)$	$O(\log V)$	$O((V + E) \log V)$
d -ary heap	$O\left(\frac{d \log V }{\log d}\right)$	$O\left(\frac{\log V }{\log d}\right)$	$O\left((V \cdot d + E) \frac{\log V }{\log d}\right)$
Fibonacci heap	$O(\log V)$	$O(1)$ (amortized)	$O(V \log V + E)$

Array implementation:

`makequeue(S)`

$A \leftarrow$ new array of values from S
return A

`decreasekey($A, i, newval$)`

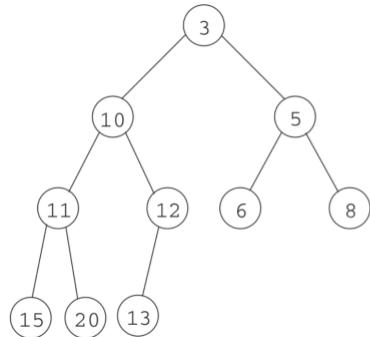
$A[i] \leftarrow newval$

`deletemin(A)`

$i \leftarrow$ index of minimum value in A
 $min \leftarrow A[i], A[i] \leftarrow \infty$
if $min = \infty$ **then return** NIL
else return min, i

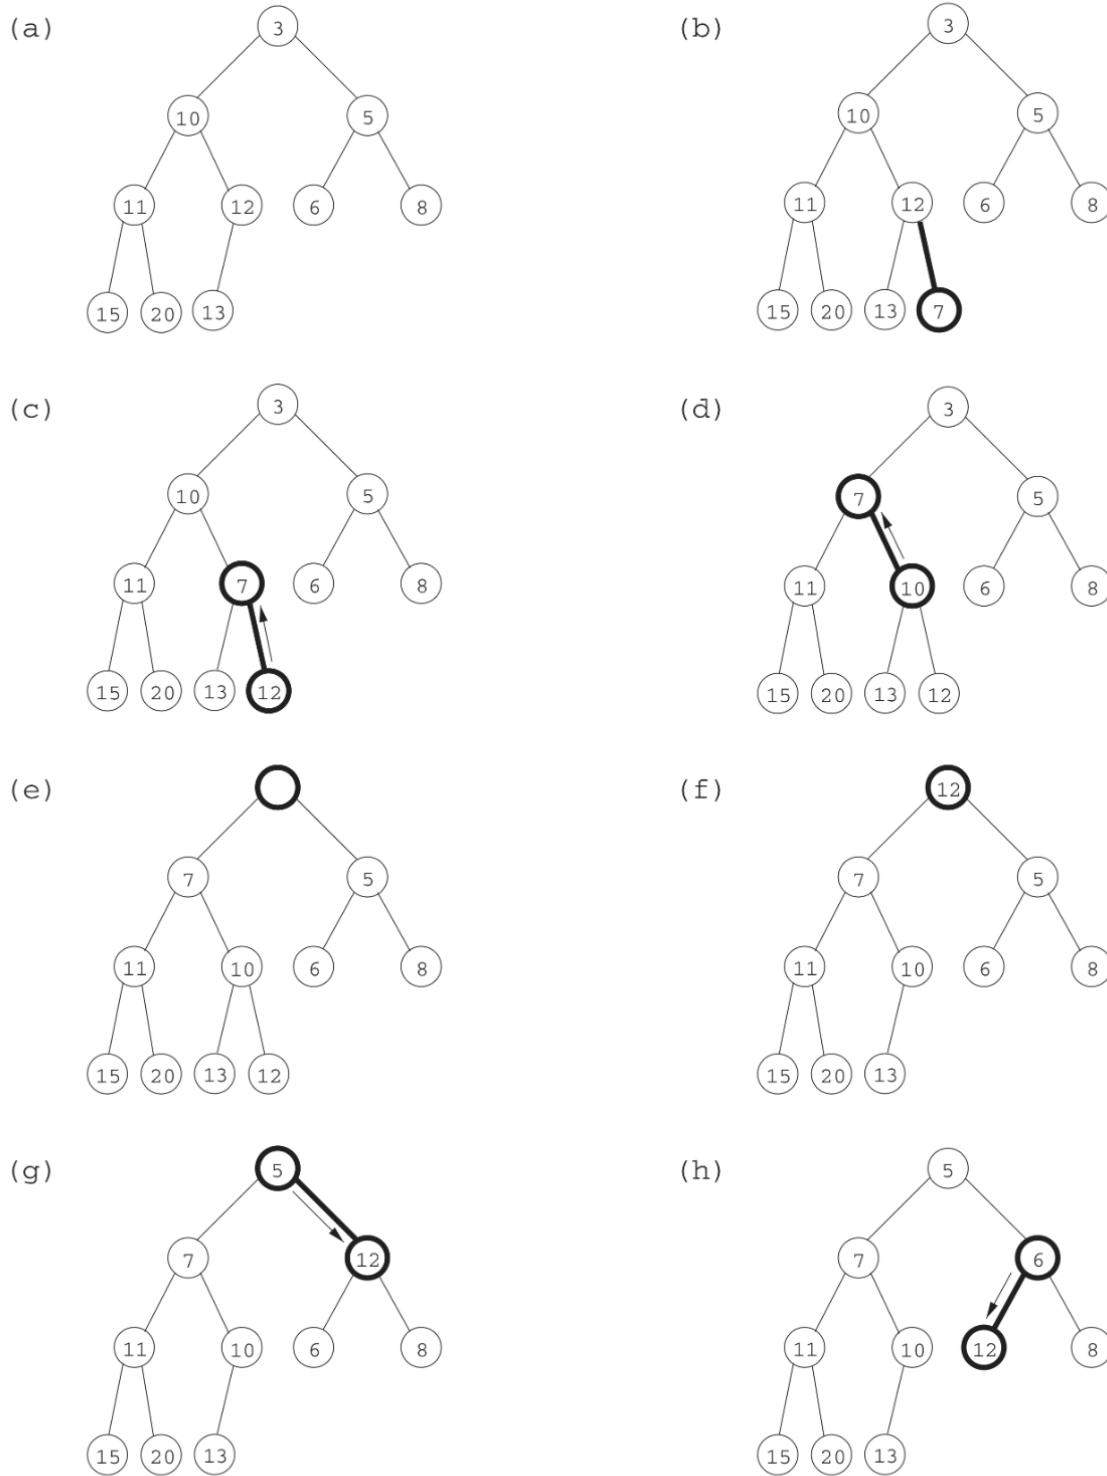
Binary heap:

1. It is a complete binary tree – every level filled from left to right and must be full before next level is started.
2. Key value of any node is \leq the key values of its children.
3. Realization as array $A[1, \dots, n]$: root at index 1, left child of a node at index i at index $2i$, right child at index $2i + 1$. Parent of node at index j is at index: $\left\lfloor \frac{j}{2} \right\rfloor$.



3	10	5	11	12	6	8	15	20	13
---	----	---	----	----	---	---	----	----	----

Figure 4.11 (a) A binary heap with 10 elements. Only the key values are shown. (b)–(d) The intermediate “bubble-up” steps in inserting an element with key 7. (e)–(g) The “sift-down” steps in a delete-min operation.



5.2 Shortest Paths and Negative Edge Weights

Dijkstra's not guaranteed to be correct. If we have a negative weight cycle reachable from source vertex, then notion of shortest path (distance) not well defined.

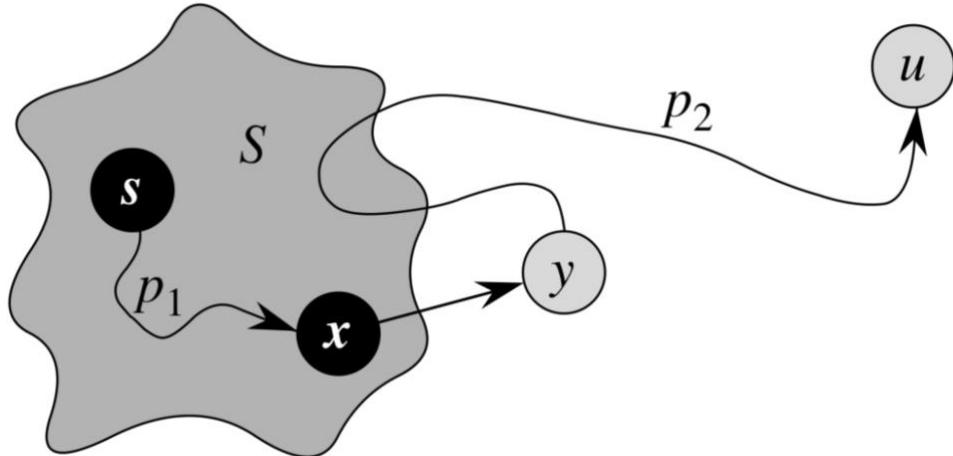
So, we assume graph has no negative weight cycles (reachable from s).

Claim 1: *In a graph that has no negative weight cycles, if there exists a path $u \rightsquigarrow v$, there exists a shortest path $u \rightsquigarrow v$ that is simple.*

Proof by construction: Suppose we are given a shortest path from u to v that contains a cycle.

Just remove the cycle and now we have a path of weight at most the original.

Where proof for correctness of Dijkstra falls apart with negative weight edges:



- Given that the $s \rightsquigarrow u$ path in the picture is a shortest path, we know that the $s \rightsquigarrow y$ subpath is a shortest path.
- But that does not imply $\delta(s, y) \leq \delta(s, u)$ unless edge weights are ≥ 0 .
 - E.g., p_2 is the single edge $y \rightarrow u$ with $l(y, u) = -1$.
 - Then, $\delta(s, u) = \delta(s, y) - 1$.

Key operation in Dijkstra's:

```
procedure update(( $u, v \in E$ )  
dist( $v$ ) = min{dist( $v$ ), dist( $u$ ) + l( $u, v$ )}
```

Claim 2: Suppose $\langle u_0, u_1, \dots, u_k \rangle$ is a shortest path $u_0 \rightsquigarrow u_k$. Then, if we initialize $\text{dist}(u_0) = 0$ and for all $i = 1, \dots, k$, $\text{dist}(u_i) = \infty$, and we perform at least one update on every edge $\langle u_i, u_{i+1} \rangle$ in that shortest path in order, then afterwards, $\text{dist}(u_k) = \delta(u_0, u_k)$.

Proof. By induction on k . Exploit Claim (2) from Lecture Notes 5(a).

5.2.1 Bellman-Ford Algorithm

Figure 4.13 The Bellman-Ford algorithm for single-source shortest paths in general graphs.

```

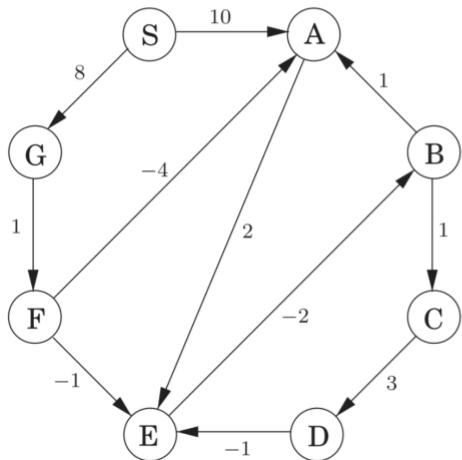
procedure shortest-paths( $G, l, s$ )
Input: Directed graph  $G = (V, E)$ ;
       edge lengths  $\{l_e : e \in E\}$  with no negative cycles;
       vertex  $s \in V$ 
Output: For all vertices  $u$  reachable from  $s$ ,  $\text{dist}(u)$  is set
       to the distance from  $s$  to  $u$ .

for all  $u \in V$ :
   $\text{dist}(u) = \infty$ 
   $\text{prev}(u) = \text{nil}$                                 max # edges in any simple path

 $\text{dist}(s) = 0$ 
repeat  $|V| - 1$  times:
  for all  $e \in E$ :
    update( $e$ )

```

Figure 4.14 The Bellman-Ford algorithm illustrated on a sample graph.



Node	Iteration							
	0	1	2	3	4	5	6	7
S	0	0	0	0	0	0	0	0
A	∞	10	10	5	5	5	5	5
B	∞	∞	∞	10	6	5	5	5
C	∞	∞	∞	∞	11	7	6	6
D	∞	∞	∞	∞	∞	14	10	9
E	∞	∞	12	8	7	7	7	7
F	∞	∞	9	9	9	9	9	9
G	∞	8	8	8	8	8	8	8

Time efficiency of Bellman-Ford: $\Theta(|V| \cdot |E|)$.

- $\Omega(|V|)$: Best case elicited by $E = \emptyset$ and $|E| = \Theta(1)$.
- $O(|V|^3)$: Worst case elicited by $|E| = \Theta(|V|^2)$.

5.2.2 Directed Acyclic Graphs (DAGs)

Shortest paths (distances) in DAGs:

After we linearize, all paths run left to right. Exploit Claim (2) above.

Figure 4.15 A single-source shortest-path algorithm for directed acyclic graphs.

```
procedure dag-shortest-paths( $G, l, s$ )
Input: Dag  $G = (V, E)$ ;
       edge lengths  $\{l_e : e \in E\}$ ; vertex  $s \in V$ 
Output: For all vertices  $u$  reachable from  $s$ ,  $\text{dist}(u)$  is set
       to the distance from  $s$  to  $u$ .

for all  $u \in V$ :
     $\text{dist}(u) = \infty$ 
     $\text{prev}(u) = \text{nil}$ 

 $\text{dist}(s) = 0$ 
Linearize  $G$ 
for each  $u \in V$ , in linearized order:
    for all edges  $(u, v) \in E$ :
        update( $u, v$ )
```
